

Bayesian Data Analysis

Bayesian Data Analysis Techniques

\neq

Applying Bayes' Theorem

Probability

- Frequentist interpretation:
Probability measures a proportion of outcomes
- Bayesian interpretation:
Probability measures a degree of belief

Usefulness

When asking

Does my code have bugs?

What is more useful:

- Yes Frequentist
- Yes (80%), No (20%) Bayesian

Goals

1. Estimate parameter values
2. Predict data values
3. Model comparison

Bayes' Theorem

Posterior

Likelihood

Prior

$$P(A|B) = \frac{P(B|A)}{P(B)} P(A)$$

Evidence

Bayes' Theorem

$$P(A|B) = \frac{P(B|A)}{\int p(B|A)p(A)dA} P(A)$$

Analytical Solution

$$P(A|B) = \frac{P(B|A)}{\int p(B|A)p(A)dA} P(A)$$

$P(A)$ conjugate prior of $P(B|A)$

Applying To Models

$p(\text{data values} | \text{model structure and parameters})$

Applying To Models

$p(\text{data values} | \text{model structure and parameters})$


For instance: coin flip

$$p(\text{coin=Heads} | \theta) = \theta$$

$$p(\text{coin=Tails} | \theta) = 1 - \theta$$

Applying To Models

$p(\text{data values} | \text{model structure and parameters})$



$p(\text{model structure and parameters} | \text{data values})$

A blue callout bubble with a dark blue outline, pointing from the text below to the second expression above.

Probability to
evaluate model

Applying To Models

$p(\text{data values} | \text{model structure and parameters})$

\Rightarrow

$p(\text{model structure and parameters} | \text{data values})$

For instance: coin flip

$p(\theta | \text{coin=Heads})$

The Scientific Method

- Most successful way till now.
 - Observe system with different inputs.
 - Guess rules.
 - Predict outputs.
 - **Check with real system!**
 - If prediction matches outputs, you **may** have the right rules.
 - If prediction doesn't match outputs, you are definitely wrong.

Updating

$$P(A|B) = \frac{P(B|A)}{P(B)} P(A)$$

Updating


$$P(A|B', B) = \frac{P(B'|A)}{P(B')} \underbrace{\frac{P(B|A)}{P(B)}}_{\text{Old posterior}} P(A)$$

Most Probable Result

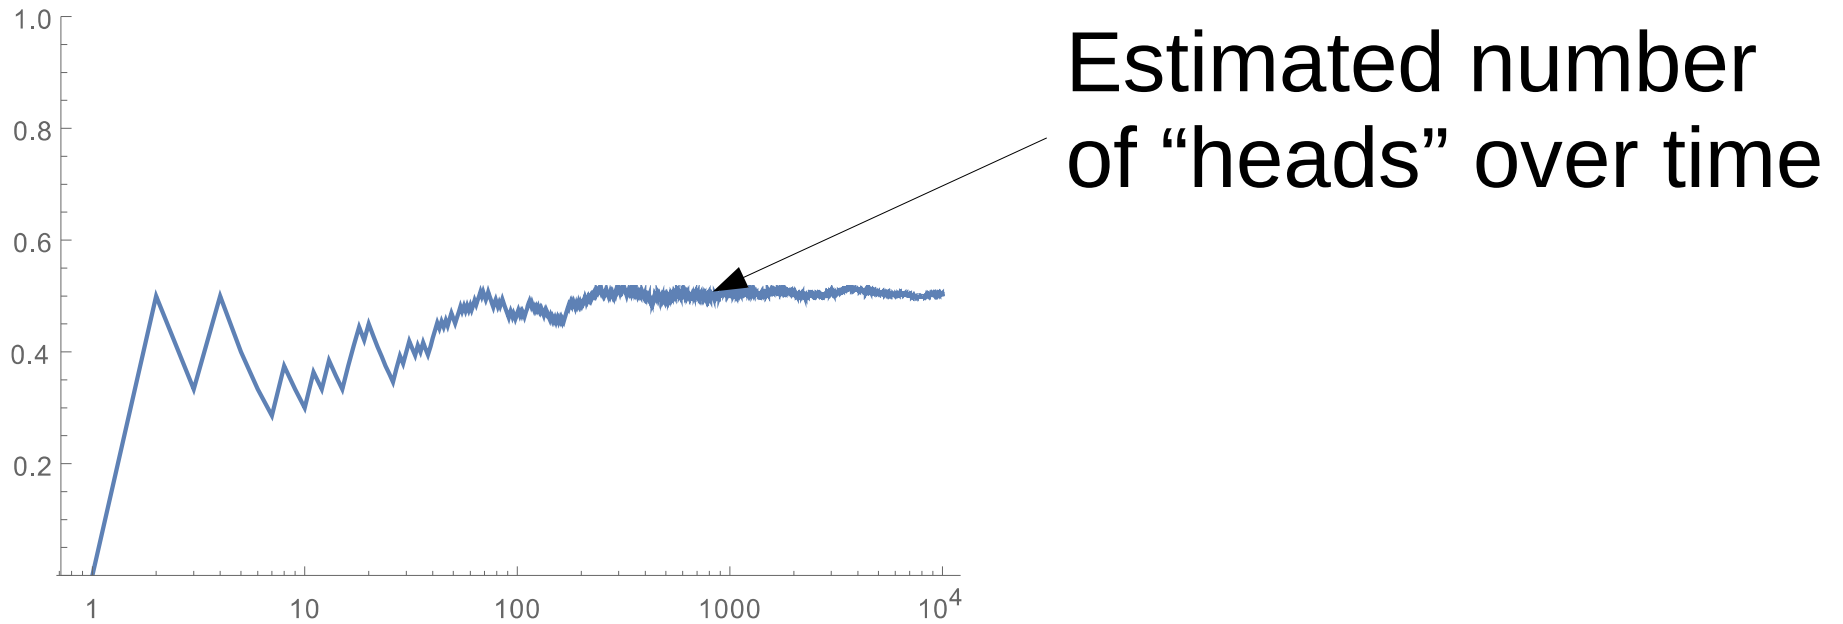
$$\hat{A} = \operatorname{argmax}_A p(A|B)$$

Coin Flips

Estimated number
of “heads” over time




Coin Flips




Reliant on "Law of Large Number"

Coin Flips

Estimated number of “heads” over time



Bayes posterior probability over time




Example 1

- Write down information about gender of 20 random people
- Draw five of the notes, three indicate female
- How many females in original group?


Example 1

- Without any knowledge: 50-50



Binomial Distribution $B(20, 0.5)$

- Maybe no knowledge



Uniform Distribution $U(0, 20)$

Prior

Example 1


- Sampling: 5 out of 20 with positive result of 3
- Hypergeometric Distribution: $H(N, K, n)$
- Need Likelihood:

Likelihood of set of parameter values θ given outcome x is equal to probability of observed outcome.

$$\mathcal{L}(20,5,3|x) = \text{PDF}(H(20,5,x), 3)$$


Example 1

- Likelihood



Example 1


- Posterior: multiply prior with likelihood and normalize



- Original Prior
- Likelihood
- Posterior

Example 1


- Sample again, with same result



- Original Prior
- Likelihood
- Posterior


Example 1

- Sample again, with same result




- Original Prior
- Likelihood
- Posterior

Do Not Be Too Sure




Use 0.01 or 0.99



Example 2: Coin Fairness

- Model for a coin:
 - Really Bernoulli with $\theta=0.5$?
 - Have some throws
 - Update model
- Beta is prior for Binomial



Example 2: Coin Fairness

- Given sequence

H, H, H, H, H, H, H, H, H, H, H, T, T, T

what is the value of θ ?

Example 2: Coin Fairness

- R code: use sampling program (JAGS)

```
modelString = "  
model {  
  # Likelihood:  
  for ( i in 1:nFlips ) {  
    y[i] ~ dbern( theta )  
  }  
  # Prior distribution:  
  theta ~ dbeta( priorA , priorB )  
}"
```



Example 2: Coin Fairness

- Initialization of parameters of JAGS program:

```
dataList = list(  
  nFlips = 14 ,  
  y = c( 1,1,1,1,1,1,1,1,1,1,1,0,0,0 ),  
  priorA = 1,  
  priorB = 1  
)
```



Example 2: Coin Fairness

JAGS Results



Even More Elaborate

- Model the mints:



Explicit Code Example

- Have data
- Have idea about model: $a \cdot x + b$

C++ Code

```
typedef std::tuple<double,double,double> chainval_t;
{
    size_t iterations = 10000;
    chainval_t startvalue = std::make_tuple(0.1, 0.5, 0.3);
    auto chain = run_metropolis_MCMC(startvalue, iterations, x, y);

    size_t burnin = 5000;
    auto accum = std::accumulate(chain.begin() + burnin, chain.end(),
                                std::make_tuple(0.0, 0.0, 0.0),
                                [](const auto& l, const auto& r){
    return std::make_tuple(std::get<0>(l) + std::get<0>(r),
                            std::get<1>(l) + std::get<1>(r),
                            std::get<2>(l) + std::get<2>(r)); });
```

C++ Code

```
auto run_metropolis_MCMC(const chainval_t& startvalue, size_t iterations, const
std::vector<double>& x,
                        const std::vector<double>& y)
{
    std::vector<chainval_t> chain;
    chain.push_back(startvalue);

    __gnu_cxx::sfmt19937 rg((std::random_device())());
    std::uniform_real_distribution<double> udf;

    while (--iterations > 0) {
        auto proposal = proposalfunction(chain.back(), rg);

        auto probab = std::exp(logposterior(proposal, x, y) - logposterior(chain.back(), x, y));
        if (udf(rg) >= probab)
            proposal = chain.back();
        chain.push_back(proposal);
    }
}
```

C++ Code

```
double logposterior(const chainval_t& param,  
                    const std::vector<double>& x,  
                    const std::vector<double>& y)  
{  
    auto a = std::get<0>(param);  
    auto b = std::get<1>(param);  
    auto sd = std::get<2>(param);  
    return loglikelihood(a, b, sd, x, y)  
        + logprior(a, b, sd);  
}
```

C++ Code

```
double loglikelihood(double a, double b, double sd,  
                    const std::vector<double>& x, const std::vector<double>& y)  
{  
    std::vector<double> sl(x.size());  
    std::transform(x.begin(), x.end(), y.begin(), sl.begin(),  
                  [a,b,sd](auto x, auto y){return dnorm<double,true>(y, a*x+b, sd);});  
    return std::accumulate(sl.begin(), sl.end(), 0.0);  
}
```

```
double logprior(double a, double b, double sd)  
{  
    auto aprior = dunif<double,true>(a, 0.0, 10.0);  
    auto bprior = dnorm<double,true>(b, 0.0, 5.0);  
    auto sdprior = dunif<double,true>(sd, 0.0, 30.0);  
    // addition because we have log probabilities  
    return aprior + bprior + sdprior;  
}
```


C++ Code

```
template<typename T, bool Log = false>
T dnorm(T x, T mu, T sd)
{
    constexpr T rt2pi = std::sqrt(2*__gnu_cxx::__math_constants<T>::__pi);
    auto a = x - mu;
    T res;
    if (Log)
        res = -a*a/(2*sd*sd)-std::log(sd)-std::log(rt2pi);
    else
        res = exp(-a*a/(2*sd*sd))/(sd * rt2pi);
    return res;
}
```

```
template<typename T, bool Log = false>
T dunif(T x, T mi, T ma)
{
    return (x < mi || x > ma
        ? (Log ? -std::numeric_limits<T>::infinity() : 0.0)
        : (Log ? std::log(1.0 / (ma - mi)) : 1.0 / (ma - mi)));
}
```

Questions?