

# Tree-based Methods

---

Sanjay Arora  
AI Center of Excellence  
Red Hat

---

Ulrich Drepper  
AI Center of Excellence  
Red Hat

# Why Decision Trees

- One of the most widely used ML algorithms
- Easy and fast to train
- Minimal data preprocessing required
- Can handle both real-valued and categorical variables
- Variants (random forests and boosted decision trees) can be extremely performant

# Problem Definition: Classification

Given features :  $x_1, x_2, \dots, x_n$

Predict class or label,  $y = 0$  or  $1$

# Problem Definition: Classification

Given features :  $x_1, x_2, \dots, x_n$

Predict class or label,  $y = 0$  or  $1$

**Will discuss regression later**

# Decision Tree

Series of if-else statements on features  
that distinguish between two classes

# Decision Tree

Series of if-else statements on features  
that distinguish between two classes


Some we code up all the time!

# Decision Tree

Series of if-else statements on features  
that distinguish between two classes


Something we code up all the time!

But now the expressions are learned!



Is it possible to separate red and green dots with a sequence of if-else statements on x and y?


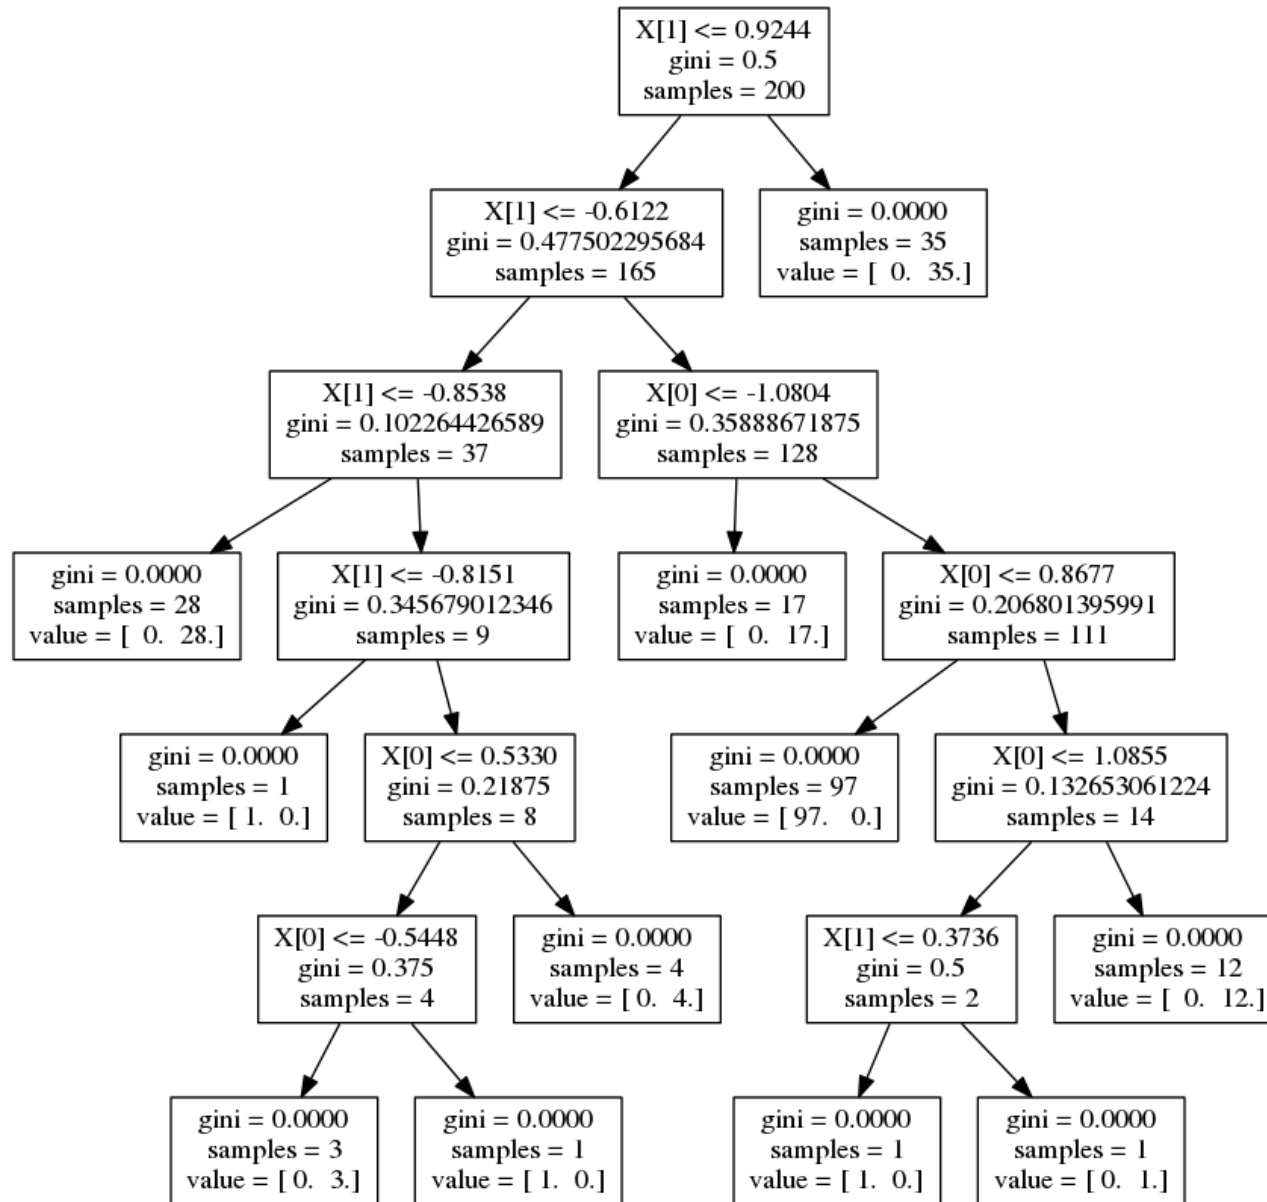


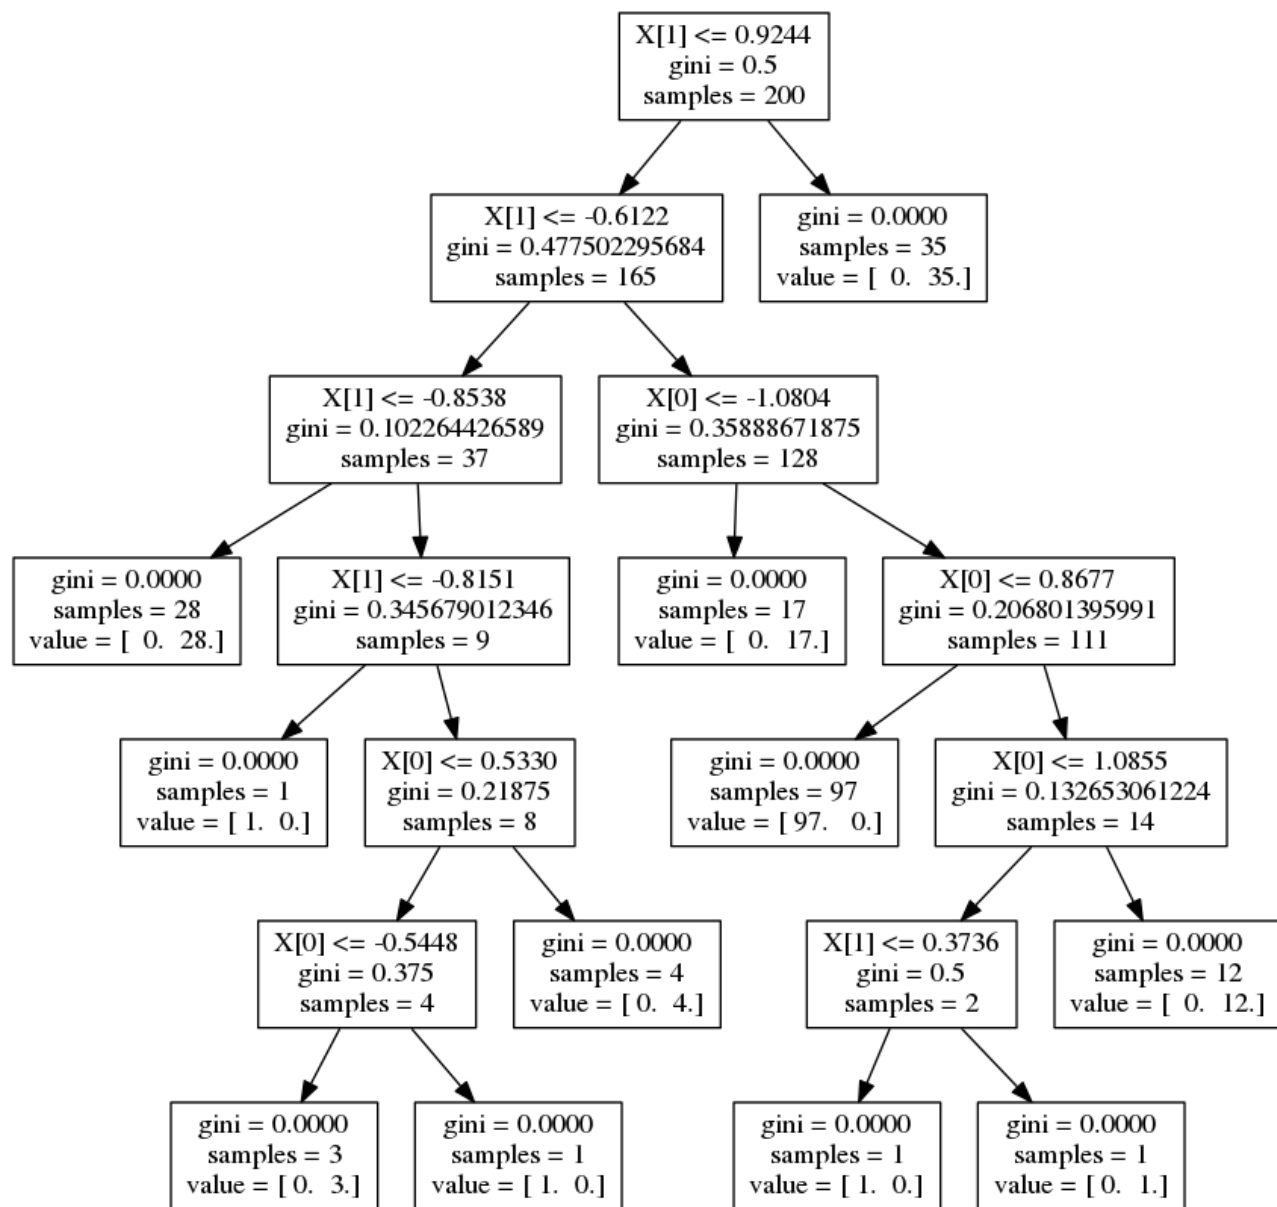


Any if-else statement on x or y





Horizontal or vertical line






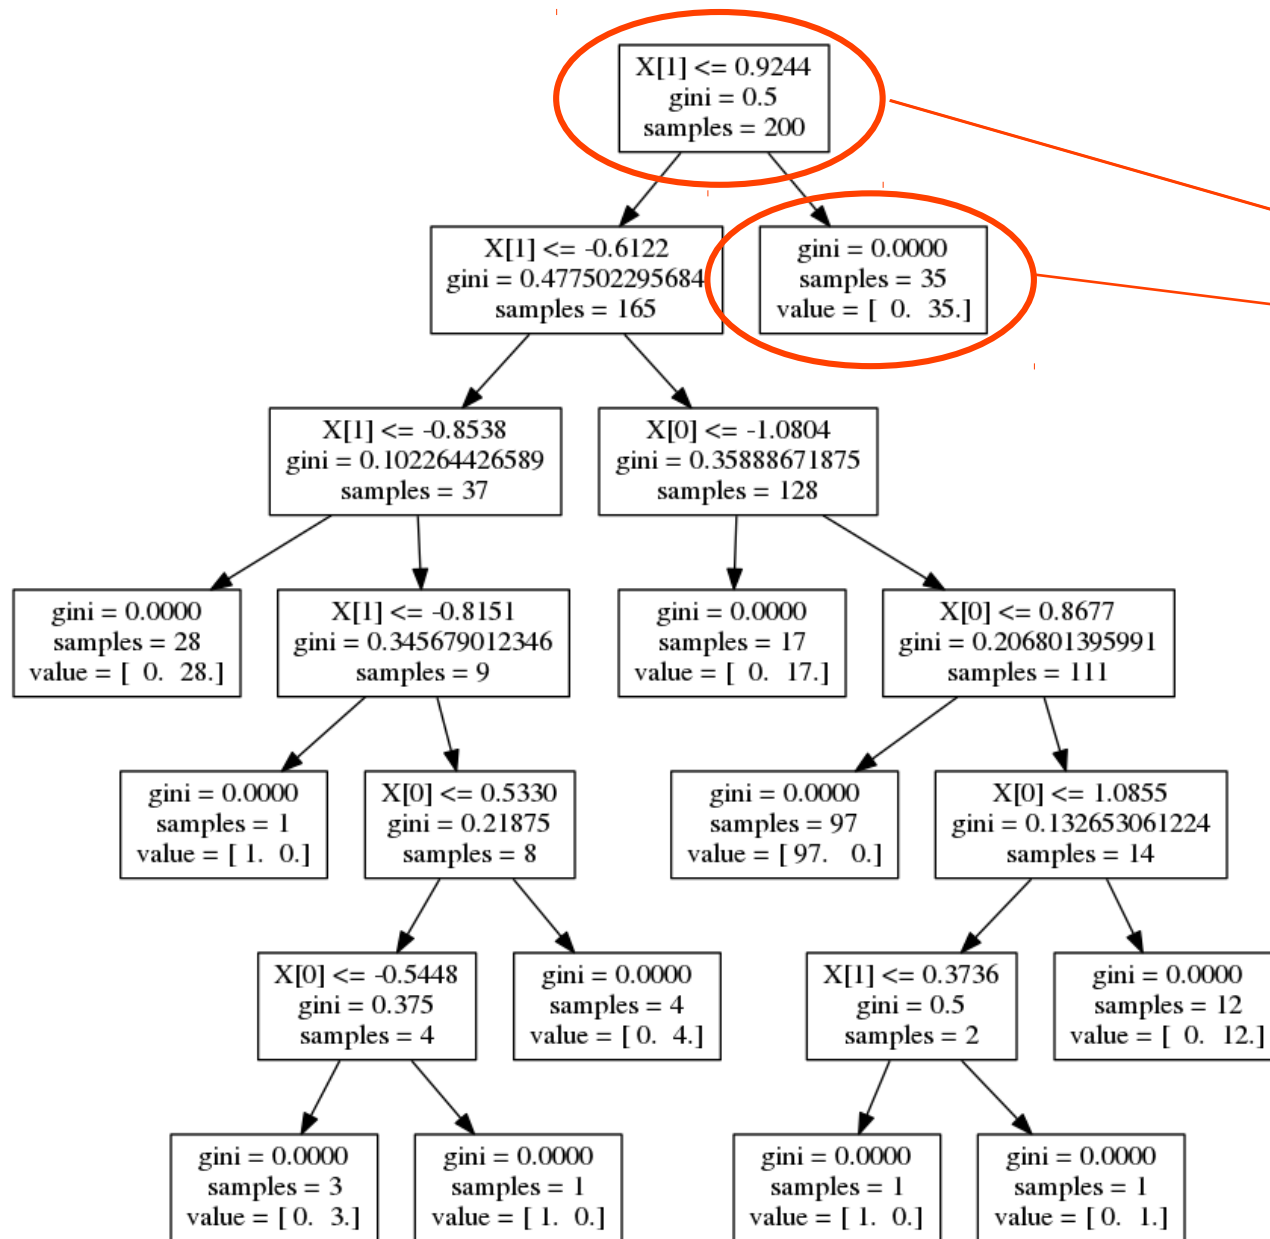
$X[0] = x$   
 $X[1] = y$






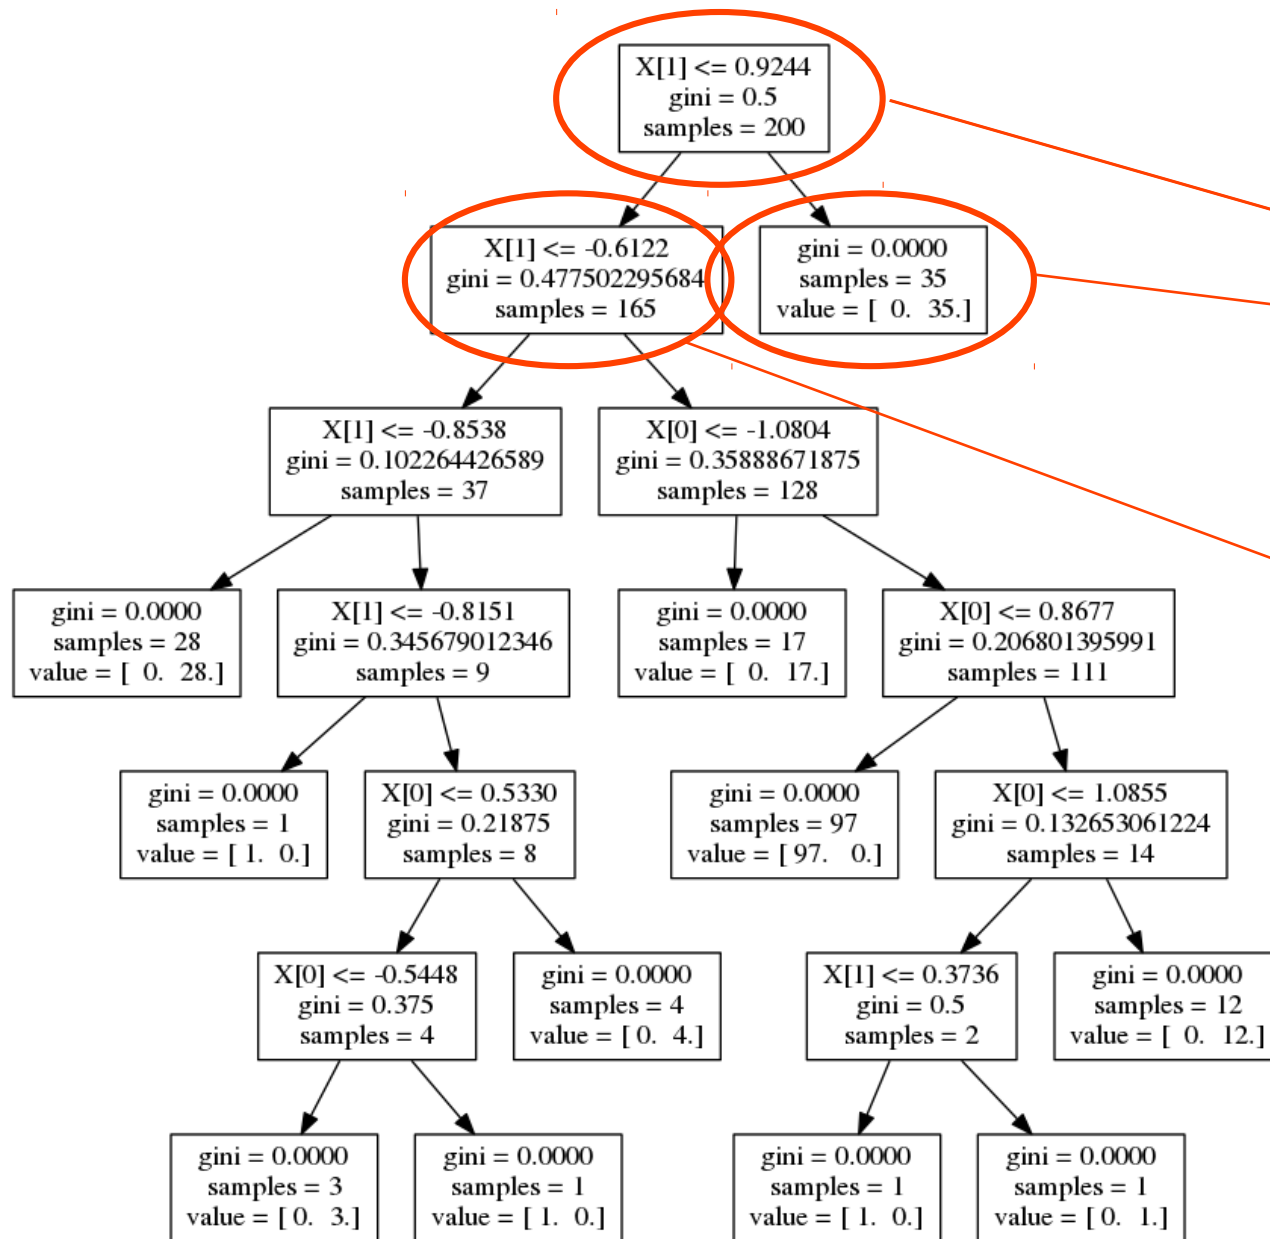
$X[0] = x$   
 $X[1] = y$






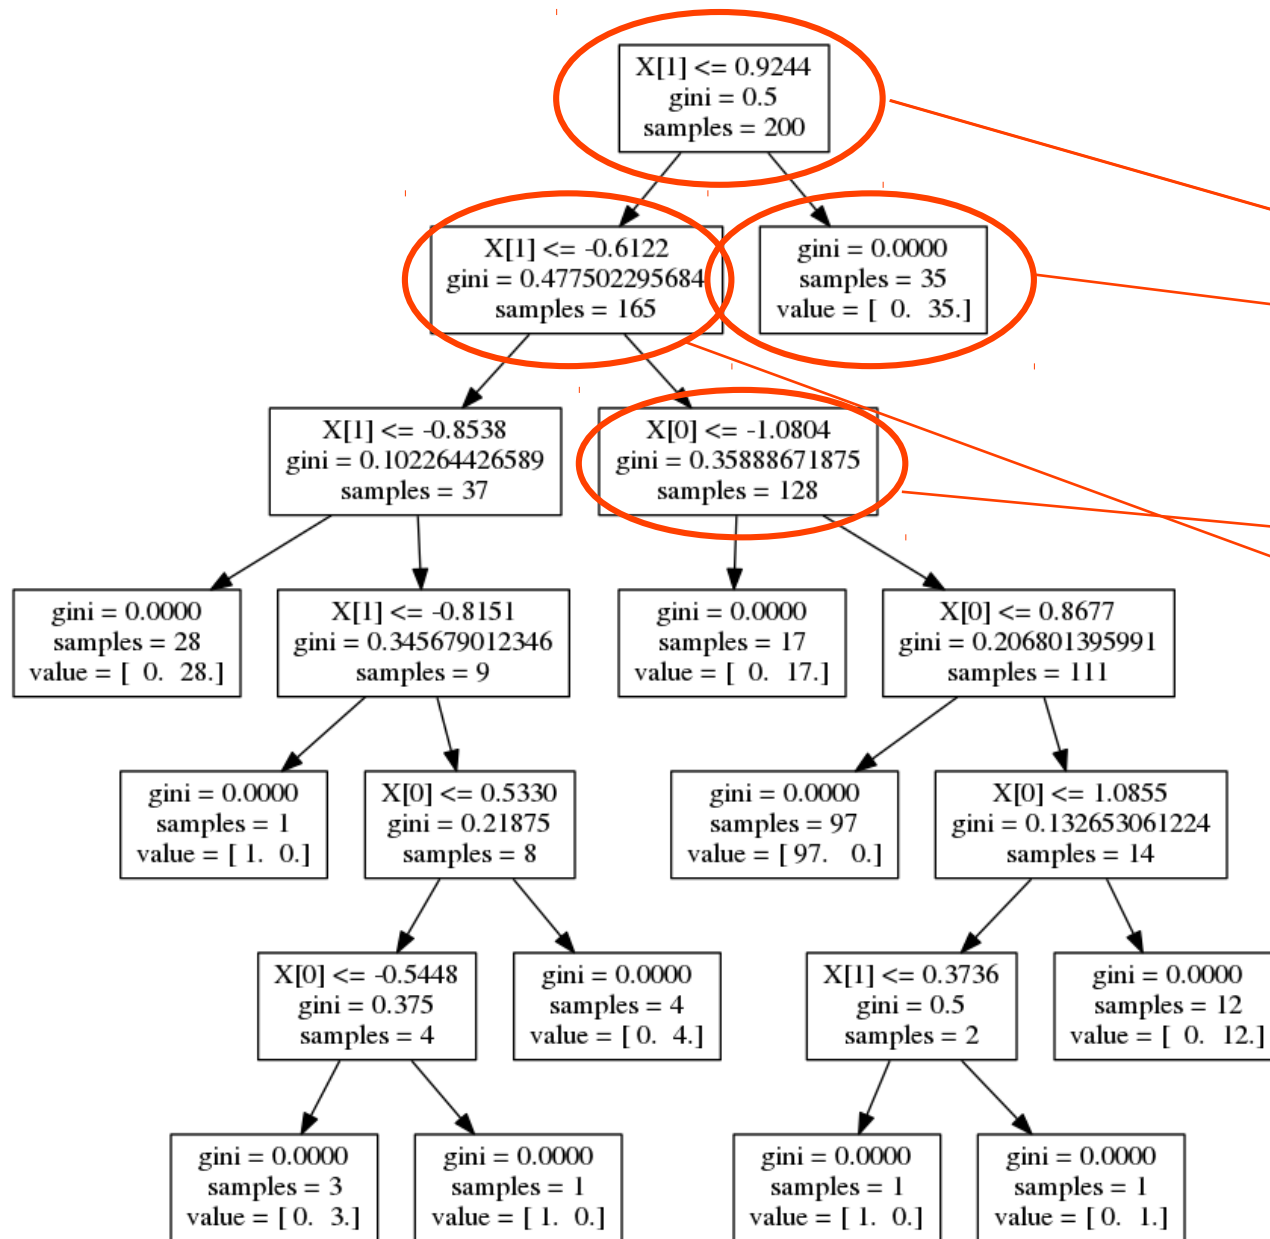
$X[0] = x$   
 $X[1] = y$






$X[0] = x$   
 $X[1] = y$





$X[0] = x$   
 $X[1] = y$



# Training

## Starting Point:

Given  $n$  features:  $x_1, \dots, x_n$

Given labels:  $y$

## Question:

What feature  $x_i$  should be choose to make a split on?

For given feature, what value should we make the split on?



# Training

## Starting Point:

Given  $n$  features:  $x_1, \dots, x_n$

Given labels:  $y$

## Question:

What feature  $x_i$  should be choose to make a split on?

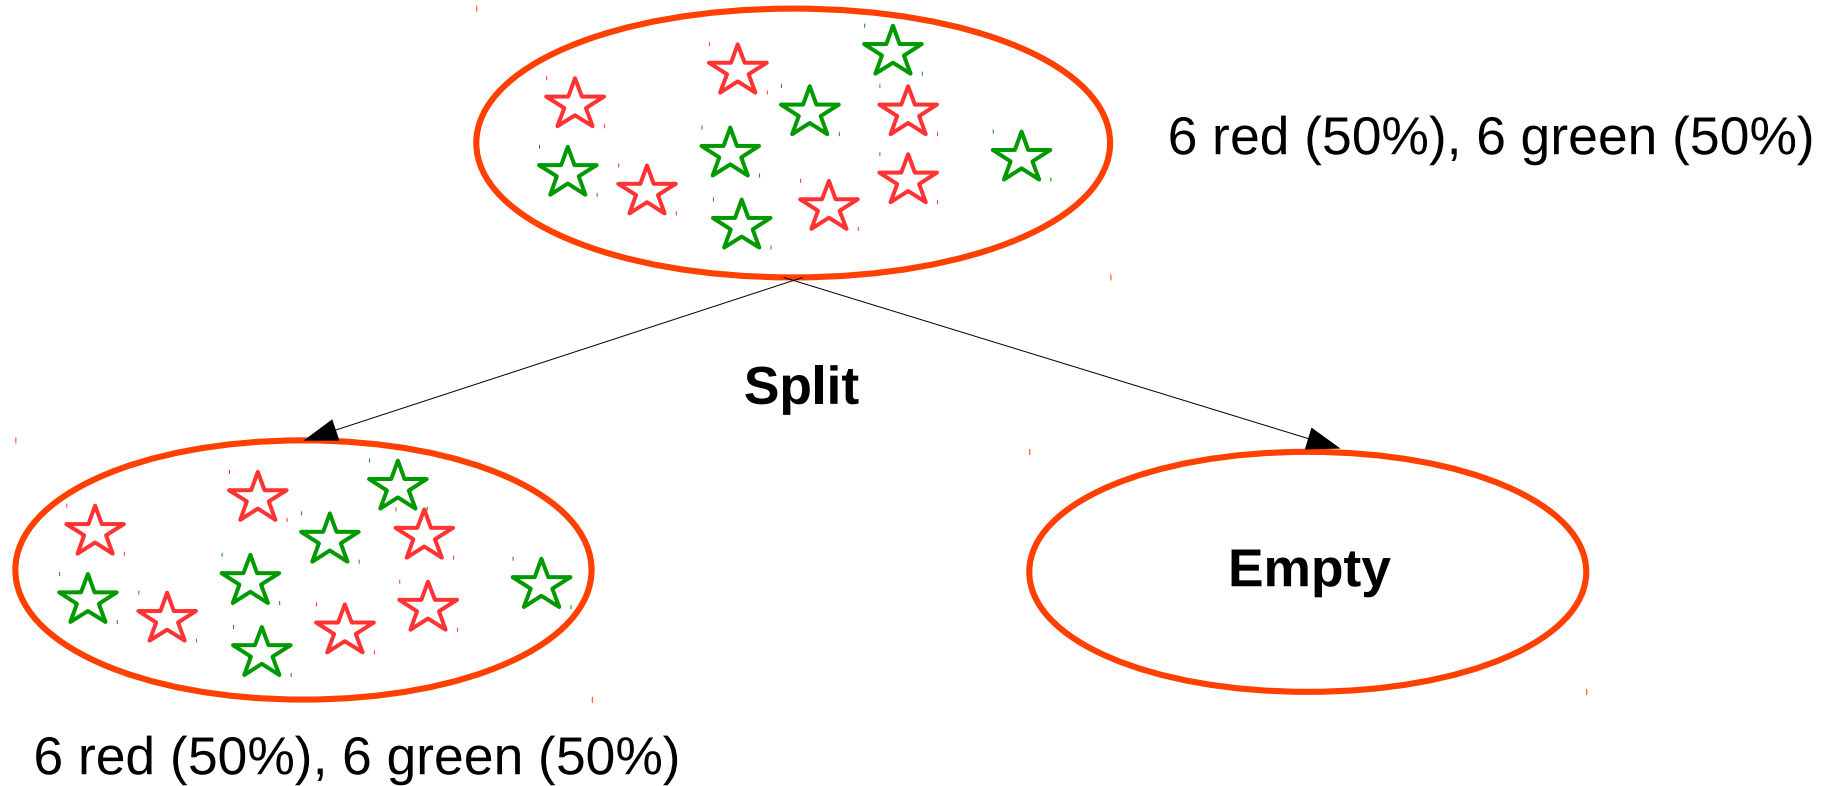
For given feature, what value should we make the split on?

## How about?:

Loop over features and all possible values to split on

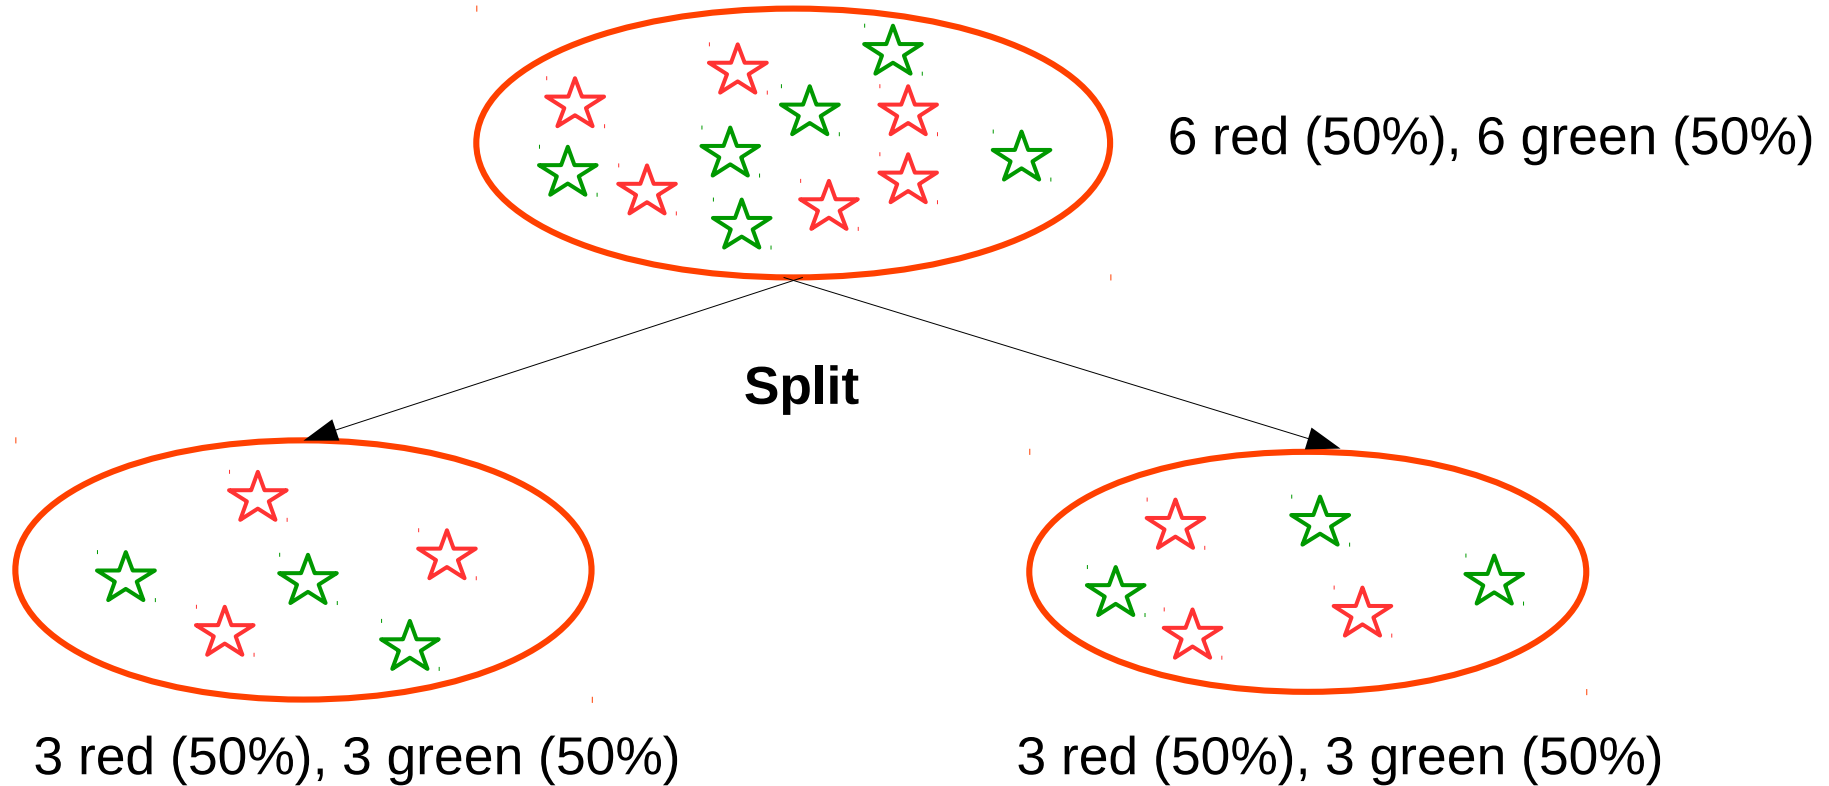
Pick one that maximizes "purity" of categories after the split

# Training




**Not useful! Still stuck with original data**

# Training






**No information gained. Still 50-50% split**

# Training



**Is this any good? Yes!! Learned a rule to separate green and red**

# Training



Split: Some function of  $p_{red}$  and  $p_{green}$  to maximize

# An Aside on Entropy

Toss a coin  $N$  times

**Result:** Get  $N_1$  heads and  $N_2$  tails

How do we describe the **state** of the system?

# An Aside on Entropy

Toss a coin  $N$  times

**Result:** Get  $N_1$  heads and  $N_2$  tails

How do we describe the **state** of the system?

**Macrostate:**  $N_1$  heads and  $N_2$  tails

# An Aside on Entropy

Toss a coin  $N$  times

**Result:** Get  $N_1$  heads and  $N_2$  tails

How do we describe the **state** of the system?

**Macrostate:**  $N_1$  heads and  $N_2$  tails

**Microstate:** Exact sequence of heads and tails

H T H H ...



# An Aside on Entropy

Toss a coin  $N$  times

**Result:** Get  $N_1$  heads and  $N_2$  tails

How do we describe the **state** of the system?

**Relevant details**

**Macrostate:**  $N_1$  heads and  $N_2$  tails

**Microstate:** Exact sequence of heads and tails

H T H H ...

**“Microscopic” description: too much detail**

# An Aside on Entropy

Define **Multiplicity**:  $\Omega$

Number of microstates corresponding to macrostate

How many detailed states correspond to one less detailed state

**Coin Toss Experiment:**  $N_1$  heads and  $N_2$  tails

$$\Omega = \binom{N_1 + N_2}{N_1} = \frac{(N_1 + N_2)!}{N_1!N_2!}$$

# An Aside on Entropy

$$\Omega = \binom{N_1 + N_2}{N_1} = \frac{(N_1 + N_2)!}{N_1! N_2!}$$

Stirling's approximation:  $N! \approx N^N e^{-N}$  for large N

$$\Omega \approx \frac{(N_1 + N_2)^{N_1 + N_2} e^{-N_1 - N_2}}{N_1^{N_1} e^{-N_1} N_2^{N_2} e^{-N_2}}$$

$$\Omega \approx \frac{(N_1 + N_2)^{N_1}}{N_1^{N_1}} \frac{(N_1 + N_2)^{N_2}}{N_2^{N_2}}$$

# An Aside on Entropy

$$\Omega \approx \frac{(N_1 + N_2)^{N_1}}{N_1^{N_1}} \frac{(N_1 + N_2)^{N_2}}{N_2^{N_2}}$$

Define:

$$p_1 = \frac{N_1}{N_1 + N_2} \quad p_2 = \frac{N_2}{N_1 + N_2}$$

$$\Omega = p_1^{-N_1} p_2^{-N_2} = p_1^{-p_1 N} p_2^{-p_2 N} = (p_1^{-p_1} p_2^{-p_2})^N$$

# An Aside on Entropy

$$\Omega = (p_1^{-p_1} p_2^{-p_2})^N$$

Define **Entropy**:

$$\begin{aligned} S &= \log \Omega = N \log (p_1^{-p_1} p_2^{-p_2}) \\ &= -N(p_1 \log p_1 + p_2 \log p_2) \end{aligned}$$

Entropy per coin

$$S = -(p_1 \log p_1 + p_2 \log p_2) \quad p_1 + p_2 = 1$$

# An Aside on Entropy

Generalize to any **discrete** distribution:

Random variables takes  $N$  values -  $a_1, \dots, a_N$   
with probabilities -  $p_1, \dots, p_N$

$$p_1 + \dots + p_N = 1$$

$$S \equiv -\sum_{i=1}^N p_i \log p_i \quad p_1 + \dots + p_N = 1$$

or **continuous** distribution:

$$S \equiv -\int p(x) \log p(x) dx \quad \int p(x) dx = 1$$

# An Aside on Entropy




$$S = -(p_1 \log p_1 + p_2 \log p_2)$$

$$p_1 = \frac{6}{12} \quad p_2 = \frac{6}{12}$$

$$S = -\frac{1}{2} \log \frac{1}{2} - \frac{1}{2} \log \frac{1}{2} = \log 2 \neq 0$$

# An Aside on Entropy




$$S = -(p_1 \log p_1 + p_2 \log p_2)$$

$$p_1 = \frac{6}{6} \qquad p_2 = \frac{0}{6}$$

$$S = -1 \log 1 - 0 \log 0 = \boxed{0} \longleftarrow \text{Pure sample}$$



# An Aside on Entropy



$$S = -(p_1 \log p_1 + p_2 \log p_2)$$

$$p_1 = \frac{6}{6} \qquad p_2 = \frac{0}{6}$$

$$S = -1 \log 1 - 0 \log 0 = \boxed{0} \leftarrow \text{Pure sample}$$

Technical Point:  $\lim_{x \rightarrow 0} x \log x \rightarrow 0$

# Training

## Starting Point:

Given  $n$  features:  $x_1, \dots, x_n$

Given labels:  $y$

## Question:

What feature  $x_i$  should be choose to make a split on?

For given feature, what value should we make the split on?


## How about?:

Loop over features and all possible values to split on

Pick one that maximizes "purity" of categories after the split

# Training

One possible choice for "purity": Entropy



Smaller entropy = purer

Minimize:

$$\frac{n_{left}}{n_{total}} S_{left} + \frac{n_{right}}{n_{total}} S_{right}$$

# Training

Minimize:

$$\frac{n_{left}}{n_{total}} S_{left} + \frac{n_{right}}{n_{total}} S_{right}$$

Repeat until:

End up with pure samples in each node

OR

Reach some maximum depth/number of levels of tree

OR

Reduction in cost is small enough

OR

Have too few samples in node

# Training

Alternatively, minimize:

$$\frac{n_{left}}{n_{total}} G_{left} + \frac{n_{right}}{n_{total}} G_{right}$$

**Gini**  $G_{node} = p_1(1 - p_1) + p_2(1 - p_2)$

0 for pure nodes

$\frac{1}{2}$  for even split

$p_1 = \frac{1}{2} \quad p_2 = \frac{1}{2}$

Repeat until:

End up with pure samples in each node

OR

Reach some maximum depth/number of levels of tree  
etc.

# Training

Pick rows

Loop over features:

    Loop over distinct values/cuts :

        Compute entropy/gini/variance if split on this feature and cut

Pick feature and cut minimizing metric

Repeat recursively for each daughter node till stopping condition

# Decision Trees: Advantages

Minimal pre-processing required - no normalization for example

No assumptions about distribution of data

Simple to interpret - if not too deep

Can discover interactions between features automatically

Run fast

# Decision Trees: Disadvantages

Use "horizontal" and "vertical" lines to learn decision boundaries

Result in over-complex trees

Can easily overfit data

Unstable: small changes in data can lead to very different trees

Cannot extrapolate




# Easy to Use

```
#Decision Tree Classifier
from sklearn import tree

model = tree.DecisionTreeClassifier()
model.fit(features, labels) #train the model

model.predict(test_features) #make predictions
#See documentation for other functions
```

# Example



Recall digits dataset

Let's use decision trees to identify the digit 3 (label = 1)  
from other digits (label = 0)

# Example

$C_i$	$pred = 0$	$pred = 1$
$label = 0$	$0 \rightarrow 0$	$0 \rightarrow 1$
$label = 1$	$1 \rightarrow 0$	$1 \rightarrow 1$

Recall confusion matrix

	pred = non-3	pred = 3
label = non-3	11052	251
label = 3	241	1056

$$\text{Accuracy} = \frac{(0 \rightarrow 0) + (1 \rightarrow 1)}{(0 \rightarrow 0) + (1 \rightarrow 1) + (0 \rightarrow 1) + (1 \rightarrow 0)} = 96.1\%$$

Ideally = 100%

Uniform Model (everything 0) = 89.7%

# Example

$C_i$	$pred = 0$	$pred = 1$
$label = 0$	$0 \rightarrow 0$	$0 \rightarrow 1$
$label = 1$	$1 \rightarrow 0$	$1 \rightarrow 1$

Recall confusion matrix

	pred = non-3	pred = 3
label = non-3	11052	251
label = 3	241	1056

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)} = 81.4\%$$

**Ideally = 100%**

**Uniform Model (everything 0) = 0%**

# Example

$C_i$	$pred = 0$	$pred = 1$
$label = 0$	$0 \rightarrow 0$	$0 \rightarrow 1$
$label = 1$	$1 \rightarrow 0$	$1 \rightarrow 1$

Recall confusion matrix

	pred = non-3	pred = 3
label = non-3	11052	251
label = 3	241	1056

$$\text{Specificity} = \frac{(0 \rightarrow 0)}{(0 \rightarrow 0) + (0 \rightarrow 1)} = 97.8\%$$

**Ideally = 100%**

**Uniform Model (everything 0) = 100%**

# Example

## Recall

Simple to interpret - if not too deep

# Example



# Decision Trees: Modifications

What if features are continuous not discrete?

Sort values for feature  $i$

Try splitting at each value in training dataset



# Decision Trees: Modifications

## Regression problems:

Minimize: 
$$\frac{n_{left}}{n_{total}} \sigma_{left}^2 + \frac{n_{right}}{n_{total}} \sigma_{right}^2$$


$\sigma_{L/R}^2$  = variance of numbers in node after split

$$= \sum_{i=1}^{N_{L/R}} \frac{(y_i - c_{L/R})^2}{N_{L/R}}$$

Annotations:

- $y_i$ : Value of  $i^{\text{th}}$  example
- $c_{L/R}$ : Mean of values in left/right node
- $N_{L/R}$ : Number of examples in left/right node

# Decision Trees: Pruning



Trees overfit easily!!!

# Decision Trees: Pruning

Trees overfit easily!!!

One possible solution:

Grow a full tree

Starting from the leaves, remove splits that don't result in large increases in entropy

# Ensembles of Trees

# Decision Trees: Bagging

What if had many trees and average results?

This gives a huge improvement in practice

# Decision Trees: Bagging

What if had many trees and average results?

This gives a huge improvement in practice

## Intuition:

Each tree performs well on a subset of data

Trees vote or average to give overall prediction

# Decision Trees: Bagging

Consider  $n$  random variables:  $X_1, \dots, X_n$   
independent and identically distributed

$$\text{mean}(X_i) \equiv \mu \qquad \text{var}(X_i) \equiv \sigma^2$$

$$\text{mean} \left( \frac{X_1 + \dots + X_n}{n} \right) = \frac{\text{mean}(X_1) + \dots + \text{mean}(X_n)}{n} = \frac{n\mu}{n} = \boxed{\mu}$$

$$\text{var} \left( \frac{X_1 + \dots + X_n}{n} \right) = \text{var} \left( \frac{1}{n} \sum_i X_i \right) = \frac{1}{n^2} \sum_i \text{var}(X_i) = \boxed{\frac{\sigma^2}{n}}$$

# Decision Trees: Bagging

Run  $n$  decision trees on overlapping subsets of dataset

**Bagging**

Average results from all  $n$  trees

If trees independent :

Averaging should give:

$$\mu \pm \frac{\sigma}{\sqrt{n}}$$

$\rightarrow \mu$  as  $n \rightarrow \infty$



# Decision Trees: Bagging

Run  $n$  decision trees on overlapping subsets of dataset

**Bagging**

Average results from all  $n$  trees

If trees independent :

Averaging should give: 
$$\mu \pm \frac{\sigma}{\sqrt{n}}$$
$$\rightarrow \mu \text{ as } n \rightarrow \infty$$

**What if trees not independent?**  
**(More precisely: errors from trees are uncorrelated)**

# Decision Trees: Bagging

Consider  $n$  random variables:  $X_1, \dots, X_n$   
**NOT** independent but identically distributed

$$\text{mean}(X_i) \equiv \mu$$

$$\text{var}(X_i) \equiv \sigma^2$$

Pairwise correlations:  $\mathbb{E}(X_i X_j) - \mathbb{E}X_i \mathbb{E}X_j \equiv \rho \sigma^2, i \neq j$

$$\text{mean} \left( \frac{X_1 + \dots + X_n}{n} \right) = \frac{\text{mean}(X_1) + \dots + \text{mean}(X_n)}{n} = \frac{n\mu}{n} = \boxed{\mu}$$

$$\text{var} \left( \frac{X_1 + \dots + X_n}{n} \right) = \boxed{\frac{(1 - \rho)\sigma^2}{n} + \rho\sigma^2}$$

# Decision Trees: Bagging

Run  $n$  decision trees on overlapping subsets of dataset

**Bagging**

Average results from all  $n$  trees

If trees not independent :

Averaging should give:

$$\mu \pm \left( \frac{(1 - \rho)\sigma^2}{n} + \rho\sigma^2 \right)$$

$$\rightarrow \mu \pm \rho\sigma^2 \text{ as } n \rightarrow \infty$$

**Correlated trees will always have variance in the average results**

# Decision Trees: Bagging

$$\text{var}\left(\frac{X_1 + \dots + X_n}{n}\right) = \frac{(1 - \rho)\sigma^2}{n} + \rho\sigma^2$$
$$\rightarrow \mu \pm \rho\sigma^2 \text{ as } n \rightarrow \infty$$

Even with many trees, have variance because of correlations

**Solution:** Reduce correlations between trees

# Reducing Correlations: Options

Train  $M$  trees

Recall training loop for a decision tree :

Pick rows

Loop over features:

    Loop over distinct values/cuts :

        Compute entropy/gini/variance if split on this feature and cut

Pick feature and cut minimizing metric

Repeat recursively for each daughter node till stopping condition

# Reducing Correlations: Options

Train  $M$  trees

Recall training loop for a decision tree :

Pick rows

Loop over features:

    Loop over distinct values/cuts :

        Compute entropy/gini/variance if split on this feature and cut

Pick feature and cut minimizing metric

Repeat recursively for each daughter node till stopping condition

# Reducing Correlations: Options

Train  $M$  trees

Recall training loop for a decision tree :

Pick rows

Loop over features:

Loop over distinct values/cuts:

Compute entropy/gini/variance if split on this feature and cut

Pick feature and cut minimizing metric

Repeat recursively for each daughter node till stopping condition

**Where can we insert some randomness?**

# Reducing Correlations: Options

Train  $M$  trees

**Option 1:** Each tree trains on fraction  $p_s$  of rows

**Option 2:** Each tree trains on fraction  $p_f$  of features/columns

**Option 3:** Each **node** only looks at a fraction  $p_n$  of features

**Option 4:** Each node only looks at a fraction  $p_e$  of features **and** cuts



# Reducing Correlations: Options

Train  $M$  trees

**Option 1:** Each tree trains on fraction  $p_s$  of rows

Pasting

# Reducing Correlations: Options

Train  $M$  trees

**Option 2:** Each tree trains on fraction  $p_f$  of features/columns

Random subspaces

# Reducing Correlations: Options

Train  $M$  trees

**Option 1:** Each tree trains on fraction  $p_s$  of rows

**Option 2:** Each tree trains on fraction  $p_f$  of features/columns

Random patches

# Reducing Correlations: Options

Train  $M$  trees

**Option 1:** Each tree trains on fraction  $p_s$  of rows  
bootstrapping (coming later)

**Option 3:** Each **node** only looks at a fraction  $p_n$  of features

Random Forest  
(used heavily)

# Reducing Correlations: Options

Train  $M$  trees

**Option 4:** Each node only looks at a fraction  $p_e$  of features **and** cuts

Extra (Extremely Randomized) Tree

# Sampling Data

Make each tree look at a subset of the data

## Bootstrapping

$N$  rows in data

Each tree draws  $N$  rows by random sampling with replacement

Some rows won't get picked

Some rows will get picked multiple times

# Sampling Data

## Bootstrapping

$N$  rows in data

Each tree draws  $N$  rows by random sampling with replacement

Each tree still trains on  $N$  rows

Rows not seen by tree form a validation dataset

These are called **Out-of-bag** or **OOB** samples

Performance of tree can be evaluated on the OOB data

# Sampling Data

## Bootstrapping

Generally used in statistics to estimate variance in measurements



# Bagging


`sklearn.ensemble.BaggingClassifier`

```
class sklearn.ensemble.BaggingClassifier(base_estimator=None, n_estimators=10,  
max_samples=1.0, max_features=1.0, bootstrap=True, bootstrap_features=False,  
oob_score=False, warm_start=False, n_jobs=None, random_state=None, verbose=0)
```

<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.BaggingClassifier.html>

# Random Forest Example

# Random Forest: Digits Example



Recall digits dataset

Let's use a random forest to identify the digit 3 (label = 1)  
from other digits (label = 0)

Using 100 trees in forest - no fine tuning of parameters

# Random Forest: Digits Example

## Decision Tree

	pred = non-3	pred = 3
label = non-3	11052	251
label = 3	241	1056

Accuracy = 96.1%

Sensitivity/Recall = 81.4%

Specificity = 97.8%

type I error = 2.2%

type II error = 18.6%

## Random Forest

	pred = non-3	pred = 3
label = non-3	11320	5
label = 3	241	1034

Accuracy = 98.0%

Sensitivity/Recall = 81.1%

Specificity = 99.9%

type I error = 0.0%

type II error = 18.9%

# Further Topics

# What we didn't cover

Boosting

Interpretation (but see practical session)

Variable Importance

Treeinterpreter - Deltas

Partial Dependence Plots

# Questions?

# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)



# Boosting: AdaBoost

# Problem Definition: Classification

Given features :  $x_1, x_2, \dots, x_n$

Predict class or label,  $y = 0$  or  $1$

Given features :  $x_1, x_2, \dots, x_n$

Predict class or label,  $y = \cancel{0}$  or 1  
 $-1$   
for convenience

# Boosting

Another way of combining various trees (or any estimators)  
to build powerful models

Different from bagging!

# AdaBoost: First Look at Boosting

What if want predictor to be of form:

$$F(x) = \text{sgn}(\underbrace{\sum_{m=1}^M \beta_m F_m(x)})$$

**Weighted sum of m predictors**


# AdaBoost: First Look at Boosting

For any supervised learning algorithm, we have a cost function:

$$C(w_0, w_1, \dots, w_n) = \sum_{i=1}^M C_i(w_0, w_1, \dots, w_n)$$

Sum over input examples

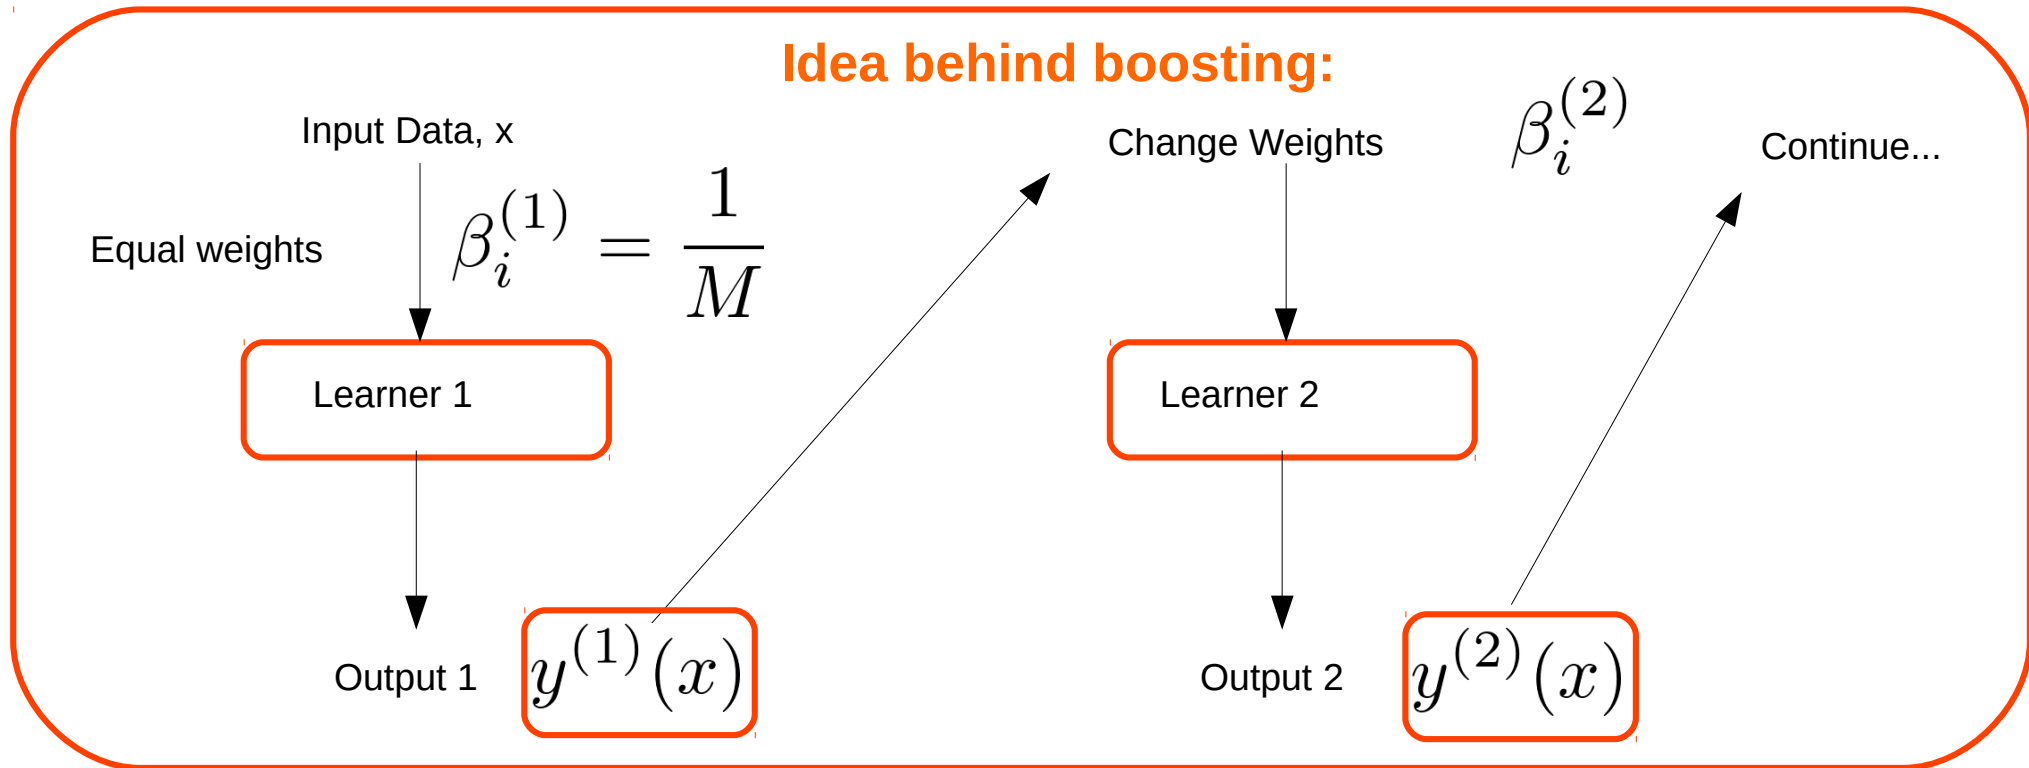
Assumption: every example **independently** and **equally** contributes to cost



$$C(w_0, w_1, \dots, w_n) = \sum_{i=1}^M \beta_i C_i(w_0, w_1, \dots, w_n)$$

Can emphasize some input examples more than others by introducing **multipliers**

# Boosting: AdaBoost



Weigh examples wrongly classified by Learner 1 more and feed to Learner 2

$$\text{Output} = \sum_{i=1}^{N_{iter}} \alpha^{(i)} y^{(i)}(x)$$

Weighing over all classifiers depending on each classifier's accuracy = **trust factor**

# Boosting: AdaBoost

Start:

$$C^{(1)}(w_0, w_1, \dots, w_n) = \sum_{i=1}^M \beta_i^{(1)} C_i(w_0, w_1, \dots, w_n)$$

$$\beta_i^{(1)} = \frac{1}{M}$$

First classifier - each example with equal weight



# Boosting: AdaBoost

Error with weights:

$$\epsilon^{(1)} = \frac{\sum_{i \in \text{Incorrect}} \beta_i^{(1)}}{\sum_{i=1}^M \beta_i^{(1)}}$$

Sum over **incorrectly predicted** examples

Sum over **all** examples

% of total weight incorrectly predicted

# Boosting: AdaBoost

Define:  $e^{\alpha^{(1)}} \equiv \frac{1 - \epsilon^{(1)}}{\epsilon^{(1)}}$

$1 - \epsilon^{(1)}$  → % of weights predicted **correctly**

$\epsilon^{(1)}$  → % of weights predicted **incorrectly**

100% accurate model →  $\epsilon^{(1)} = 0$  →  $e^{\alpha^{(1)}} = \infty$  →  $\alpha^{(1)} = \infty$

50% accurate model →  $\epsilon^{(1)} = 0.50$  →  $e^{\alpha^{(1)}} = 1$  →  $\alpha^{(1)} = 0$

0% accurate model →  $\epsilon^{(1)} = 1$  →  $e^{\alpha^{(1)}} = 0$  →  $\alpha^{(1)} = -\infty$

# Boosting: AdaBoost

Update weights:

$\beta_i^{(2)} = \beta_i^{(1)}$  if example  $i$  correctly predicted

$\beta_i^{(2)} = \beta_i^{(1)} \underbrace{e^{\alpha^{(1)}}}_{\frac{1 - \epsilon^{(1)}}{\epsilon^{(1)}}}$  if example  $i$  incorrectly predicted

# Boosting: AdaBoost

General algorithm: Learner  $m \rightarrow$  Learner  $m+1$

Minimize:  $C^{(m)}(w_0, w_1, \dots, w_n) = \sum_{i=1}^M \beta_i^{(m)} C_i(w_0, w_1, \dots, w_n)$

Compute:  $\epsilon^{(m)} = \frac{\sum_{i \in \text{Incorrect}} \beta_i^{(m)}}{\sum_{i=1}^M \beta_i^{(m)}} \quad e^{\alpha^{(m)}} \equiv \frac{1 - \epsilon^{(m)}}{\epsilon^{(m)}}$

Update:  $\beta_i^{(m+1)} = \beta_i^{(m)}$  if example  $i$  correctly classified

$\beta_i^{(m+1)} = \beta_i^{(m)} e^{\alpha^{(m)}}$  if example  $i$  incorrectly classified