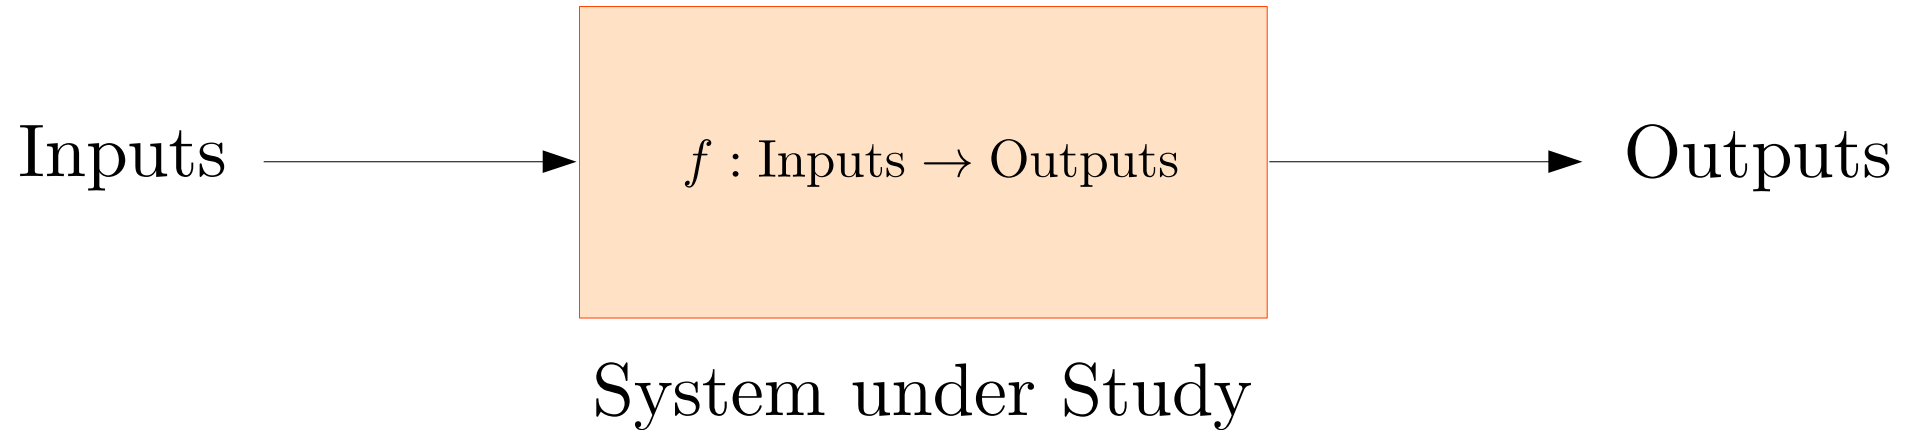


What is Machine Learning?

Sanjay Arora
AI Center of Excellence
Red Hat

Ulrich Drepper
AI Center of Excellence
Red Hat

Why and how does something work?



Question: what are the rules governing the system i.e. what is f ?

The Scientific Method

Most successful approach till now:

- Observe system with different inputs and measure outputs
- Guess rules or guess f
- Make predictions from guessed rules
- Compare predictions to outputs
- If predictions “match” outputs, you **might** have the correct rules
- If predictions don’t match outputs, you are definitely wrong

The Scientific Method

- A scientific theory can never be **proven** to be correct. There's no proof the sun will rise tomorrow – just an overwhelming possibility.
- Comparison of the predictions and outputs is more subtle than it appears. It depends on the degree of accuracy required. Rules might predict outputs within 10% but not within 1% and that might be sufficient to explain the core elements of the system under study.
- Simplicity is key. The best (in terms of predictive power) scientific theories tend to be simple.

How to guess the rules?

Many ways – see Einstein, Dirac, Feynman's work for some examples

One way:

Collect data and scan for patterns. How Kepler found the laws for planetary motion

What is Machine Learning?

- Collect data and find patterns to discover rules.
- Let computers scan through several guesses/hypotheses and find the best ones.
- Need to invent algorithms that can scan multiple hypotheses efficiently (in time, memory, amount of data required).

Terminology: Model

Our guess of the rules

Terminology: Learning

Process of scanning data to find the best rules

Terminology: Supervised Learning

Learning when data has labels (feedback) i.e.
we have both the inputs and the corresponding
outputs

Terminology: Unsupervised Learning

Learning when data doesn't have labels (no feedback)

Very common in real-life AND much
harder than supervised learning

Terminology: Reinforcement Learning

Learning when feedback infrequent in time

Closely related to control theory

Models dynamics of an agent interacting with
a complex dynamic environment

Some Lingo

Data:

Table where each column is a feature and each row is the result of one experiment.

Some Lingo

Features:

Numbers/categories collected that can potentially affect the final result. Usually, the input data has one column for each feature.

Some Lingo

Results/Labels/Targets:

Quantity to predict: Can be a continuous number or a discrete category (0/1/2 etc.).

Some Lingo

Training data:

A randomly selected percentage (say 70%) of the rows of the data.
Rest of data ignored for model-building

Some Lingo

Test data:

Part of data not included in training data (say 30%). The model built using training data is applied to test data to see how well it works on an independent set.

Some Lingo

Cleaning data:

The procedure involving filling missing data, doing sanity checks on ranges and values of various features. Tends to be an ad-hoc process.

Some Lingo

Supervised Learning:

Building a model where for each input, an output value is known.

Examples: Predict email is spam or non-spam based on its contents.
Predict height of a person from weight.

Some Lingo

Supervised Learning - **Regression:**

A supervised learning model where the output value being predicted is a real number.

Examples: Predict height of a person from weight.

Some Lingo

Supervised Learning - **Classification:**

A supervised learning model where the output value being predicted is a category.

Examples: Predict email is spam or non-spam based on its contents.

Some Lingo

Unsupervised Learning:

Build a model where nothing is being predicted. Techniques used to find patterns within data.

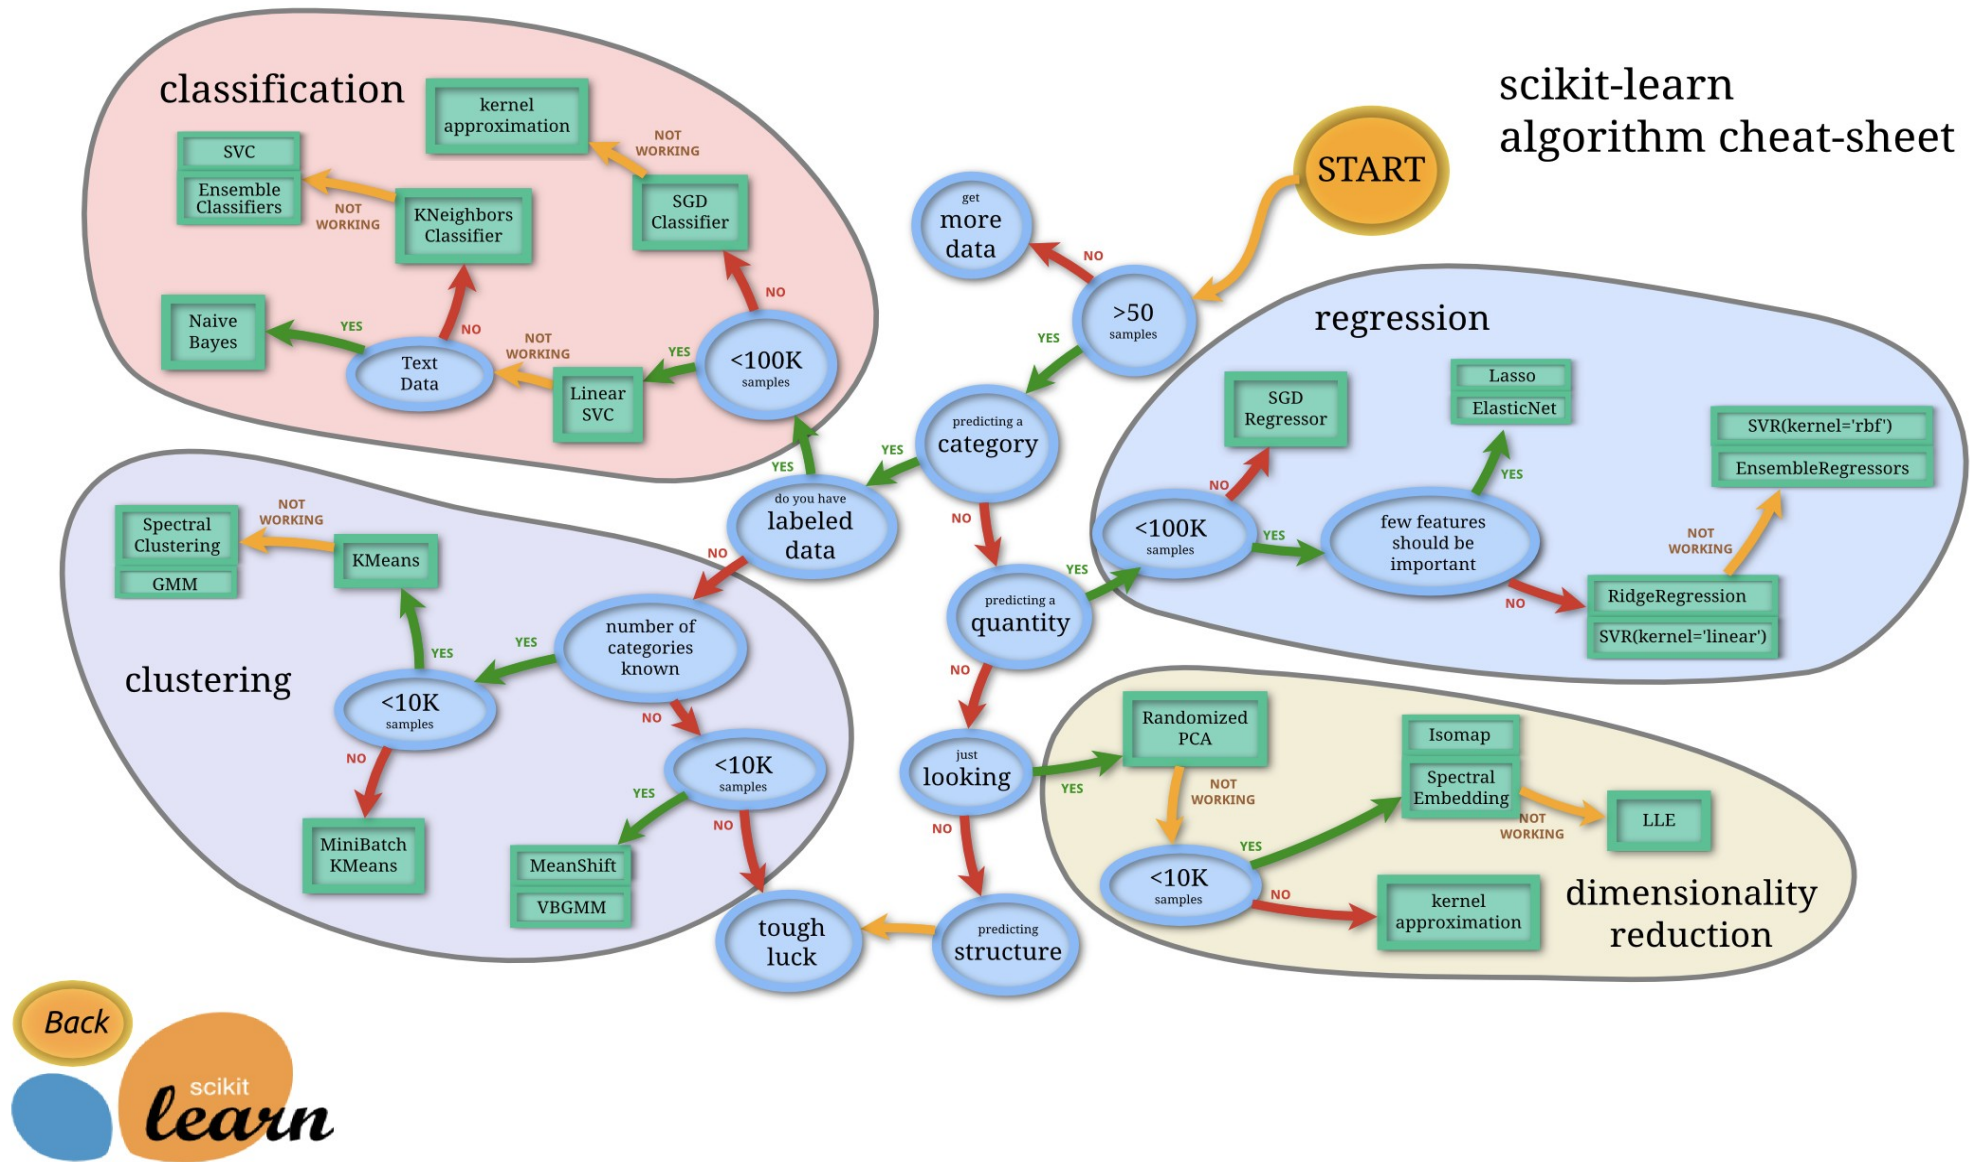
Example: clustering/grouping of data based on their pairwise distances.

Survey of Techniques: Regression

Problem Statement

Given features : x_1, x_2, \dots, x_n

predict real number y



Don't take this too literally!

Linear Regression

$$y = w_0 + w_1x_1 + \dots + w_nx_n$$

Linear in weights, **NOT** features

For example, maybe $x_1 = x$ and $x_n = x^n, n > 1$
Definitely not linear in features x , but linear in weights w_i

Linear Regression: Assumptions

Data: $y = w_0 + w_1x_1 + \dots + w_nx_n + \mathcal{N}(0, \sigma^2)$

Gaussian noise


Linear relationship between y and $x_m, \forall m$
where all other x_n are fixed

$\mathcal{N}(0, \sigma^2)$: **Gaussian** noise with **0 mean** and **constant variance**

Might consist of many noise terms

They should be **independent**, **Gaussian** with
constant variance

Example



$$y = w_0 + w_1x + w_2x^2 = 0.344 + 0.551x - 0.031x^2$$

Linear Regression: Preprocessing

Generally, features x_m can have different ranges they take values in

Good to normalized each feature to have mean = 0 and standard deviation = 1

$$x_m \rightarrow \frac{x_m - \mu(x_m)}{\sigma(x_m)}$$

Mean across all examples

Standard dev. across all examples

Mean = 0 helps with convergence of algorithms too

If features in similar ranges, can compare weights w_m to rank features according to influence on output

Linear Regression: Python Code

```
from sklearn import linear_model

model = linear_model.LinearRegression() #create instance of linear regression model
model.fit(x_values, y_values) #train model

#DONE!!!

model.predict(new_x_values) #predict result on new inputs

model.coef_ #see weights

model.score(x_values, y_values) #we'll see this soon. Compute R2
```

Linear Regression: Python Code

```
model = linear_model.Ridge(alpha = 0.5) #"Ridge" regression with one free parameter  
model = linear_model.Lasso(alpha = 0.5) #"Lasso" regression with one free parameter
```

Use these instead of linear model used previously
Prevent overfitting (coming up soon)

Also, it's very enlightening to look at source code if interested

Linear Regression: Performance

$$\text{Residual} \equiv \sum_{i=1}^N (y_i - f(x_i))^2$$

Result of
example i

Model prediction for
example i

$$\text{Variance} \equiv \sum_{i=1}^N (y_i - \bar{y})^2$$

$$\text{Unexplained Variance Ratio} = \frac{\text{Residual}}{\text{Variance}}$$


$$R^2 \equiv 1 - \frac{\text{Residual}}{\text{Variance}}$$

Close to 1 = Good(*)

Linear Regression: Performance

$$R^2 \equiv 1 - \frac{\text{Residual}}{\text{Variance}}$$

Close to 1 but might be overfitting



Compare R^2 between train and test sets!!!!

Linear Regression: Performance

Look at distribution of residuals: $y_i - f(x_i)$

Should be Gaussian with constant variance

K-Nearest Neighbors

Find K "closest" examples

Average their results

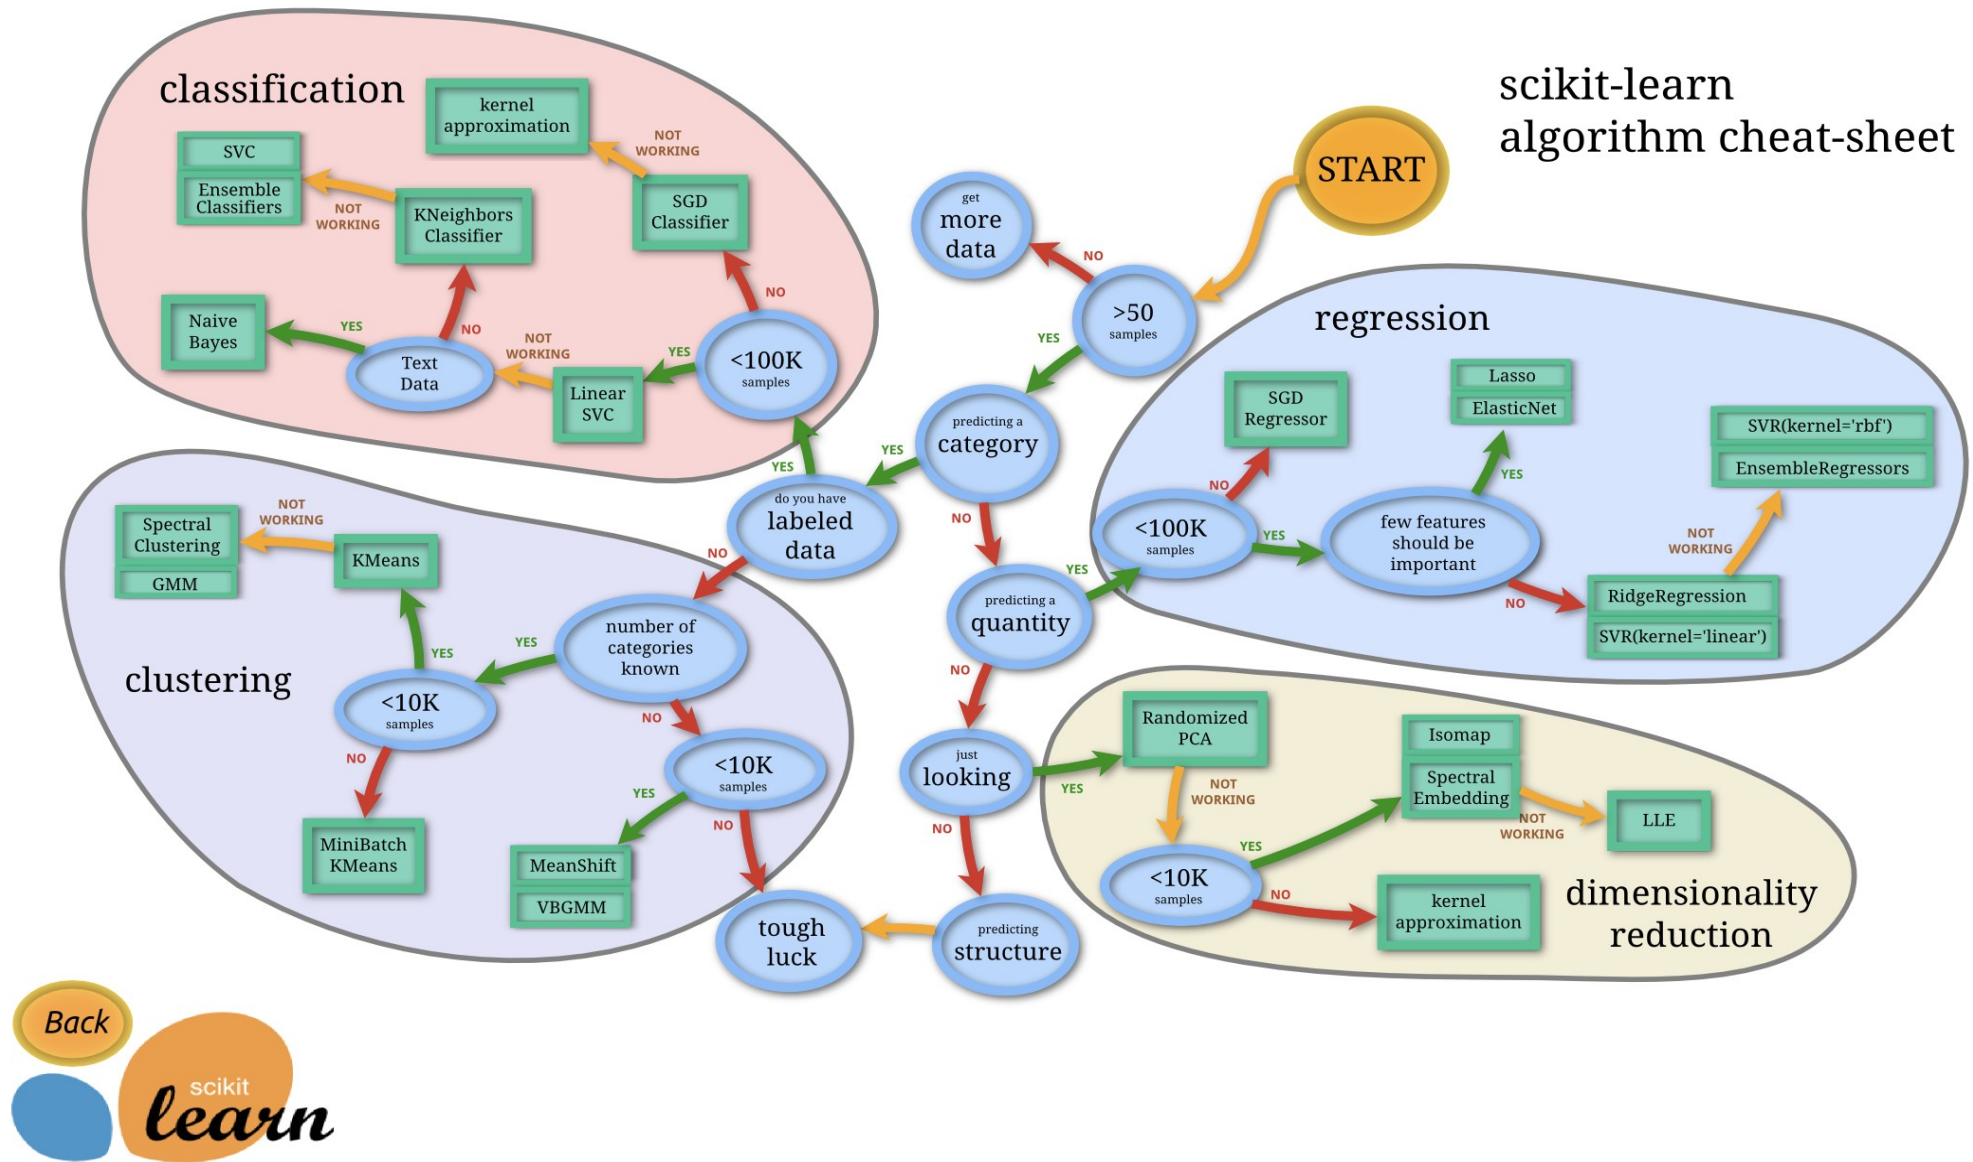
Really! That's it

Survey of Techniques: Classification

Problem Statement

Given features : x_1, x_2, \dots, x_n

predict class or label, $y = 0$ or 1



Logistic “Regression”

How about:

$$(x_1, x_2, \dots, x_n)$$



$$w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

Linear combination:



$$[0, 1]$$



Interpret as probability of
label = 1

Logistic/Sigmoid Function

Want to map real number to $[0,1]$

One possible solution: $f(x) = \frac{1}{1 + e^{-x}}$

$$x \rightarrow \infty, f(x) \rightarrow \frac{1}{1 + 0} = 1$$


$$x \rightarrow -\infty, f(x) \rightarrow \frac{1}{1 + \infty} = 0$$

$$x = 0, f(x) = 0.5$$

Logistic/Sigmoid Function

$$x \rightarrow -\infty, f(x) \rightarrow \frac{1}{1 + \infty} = 0$$

$$x \rightarrow \infty, f(x) \rightarrow \frac{1}{1 + 0} = 1$$



Probability

Want to map real number to $[0,1]$

One possible solution: $f(x) = \frac{1}{1 + e^{-x}}$

$$x = 0, f(x) = 0.5 \longrightarrow \text{"Boundary" or "Threshold"}$$

$$(x_1, x_2, \dots, x_n)$$



$$w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

Linear combination:




$$f(w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n) = \frac{1}{1 + e^{-(w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n)}}$$

< 0.5

= 0.5

> 0.5




Label 0

Boundary separating classes:

Label 1


$$w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n = 0$$

Example




Decision Boundary : $-4.86 + 1.05x + 1.00y = 0$

Example




Example



Bad!!!

Decision Boundary: $w_0 + w_1x + w_2y = 0 \implies$ Straight Line

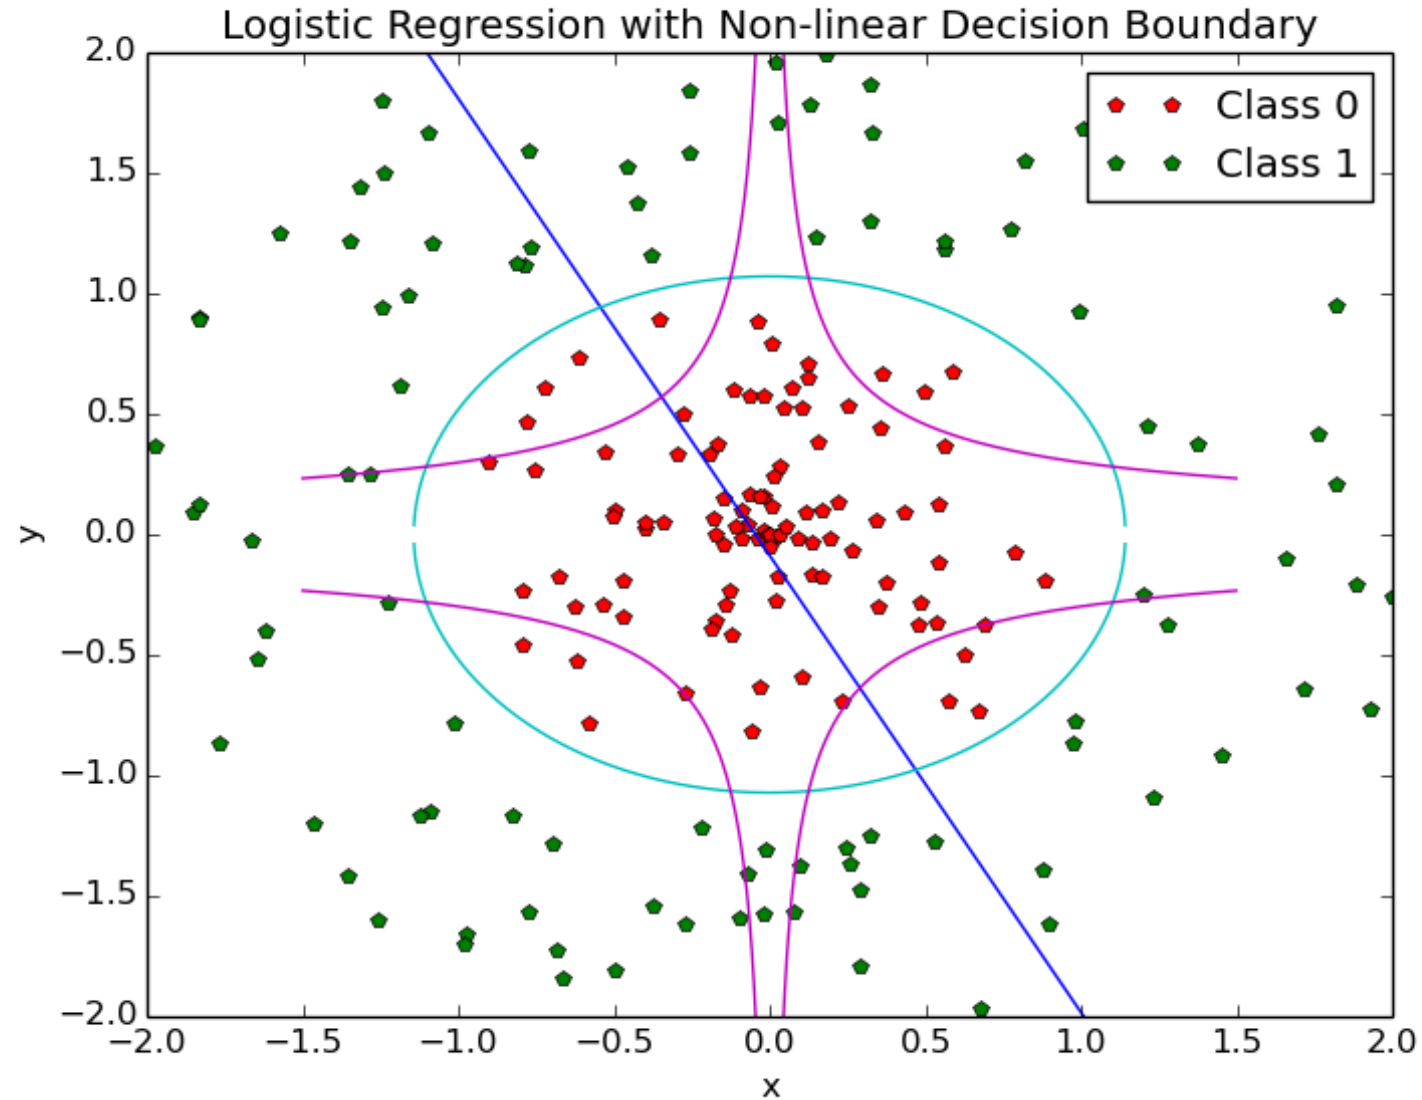
Features are: (x^2, y^2) NOT (x, y)



Accuracy = 100%

Decision Boundary: $w_0 + w_1x^2 + w_2y^2 = 0 \implies$ Ellipse (in this case at least)

Features are: $(\log |x|, \log |y|)$ NOT (x, y)



Decision Boundary: $a \log |x| + b \log |y| + c = 0 \implies y = \pm \frac{e^{-\frac{c}{b}}}{|x|^{\frac{a}{b}}}$ Red Hat

Logistic Regression: Python Code

```
from sklearn import linear_model

model = linear_model.LogisticRegression()
model.fit(features, labels) #train the model

model.predict(test_features) #make predictions (0/1)
model.predict_proba(test_features) #predict probabilities
```

Metrics for evaluation: Classification

These metrics apply to **any** classification model's output

C_i	$pred = 0$	$pred = 1$
$label = 0$	$0 \rightarrow 0$	$0 \rightarrow 1$
$label = 1$	$1 \rightarrow 0$	$1 \rightarrow 1$

$$i \rightarrow j = \text{label } i \rightarrow \text{pred } j$$

C_i	$pred = 0$	$pred = 1$
$label = 0$	True Negatives	False Positives
$label = 1$	False Negatives	True Positives

Metrics for evaluation: Classification

Good

C_i	$pred = 0$	$pred = 1$
$label = 0$	$0 \rightarrow 0$	$0 \rightarrow 1$
$label = 1$	$1 \rightarrow 0$	$1 \rightarrow 1$

Bad

$$i \rightarrow j = \text{label } i \rightarrow \text{pred } j$$

C_i	$pred = 0$	$pred = 1$
$label = 0$	True Negatives	False Positives
$label = 1$	False Negatives	True Positives

Metrics for evaluation: Classification

$$\text{Accuracy} = \frac{\# \text{ correct}}{\# \text{ total}} = \frac{(0 \rightarrow 0) + (1 \rightarrow 1)}{(0 \rightarrow 0) + (1 \rightarrow 1) + (0 \rightarrow 1) + (1 \rightarrow 0)}$$

What if 98% of examples are labeled 0 and 2% labeled 1?

Build "dumb" model that predicts 0 for every examples

Accuracy = 98% (Get everything labeled 0 correct)

Bad measure if population of labeled classes very uneven

Metrics for evaluation: Classification

Track metrics:

$$\text{Accuracy} = \frac{(0 \rightarrow 0) + (1 \rightarrow 1)}{(0 \rightarrow 0) + (1 \rightarrow 1) + (0 \rightarrow 1) + (1 \rightarrow 0)} \quad \text{Ideal Case: 100\%}$$

$$\text{Specificity} = \frac{(0 \rightarrow 0)}{(0 \rightarrow 0) + (0 \rightarrow 1)}$$

Ideal Case: 100%

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)}$$

Ideal Case: 100%

$$\text{Precision} = \frac{(1 \rightarrow 1)}{(0 \rightarrow 1) + (1 \rightarrow 1)} \quad \text{Ideal Case: 100\%}$$

$$\text{type I error} = \frac{(0 \rightarrow 1)}{(0 \rightarrow 1) + (0 \rightarrow 0)}$$

Ideal Case: 0%

$$\text{type II error} = \frac{(1 \rightarrow 0)}{(1 \rightarrow 0) + (1 \rightarrow 1)}$$

Ideal Case: 0%

Precision and Recall

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)}$$

% of things that are labeled 1
were predicted to be 1

*e.g. in a labeled dataset, if there are 200
fraudulent credit card transactions, what %
are predicted to be fraudulent by model*

$$\text{Precision} = \frac{(1 \rightarrow 1)}{(0 \rightarrow 1) + (1 \rightarrow 1)}$$

Out of things that the model predicts should
be 1, what % are actually 1s

*e.g. if a model predicts 100 credit card
transactions are fraud, what % are actually
frauds - relevant to the end-user*

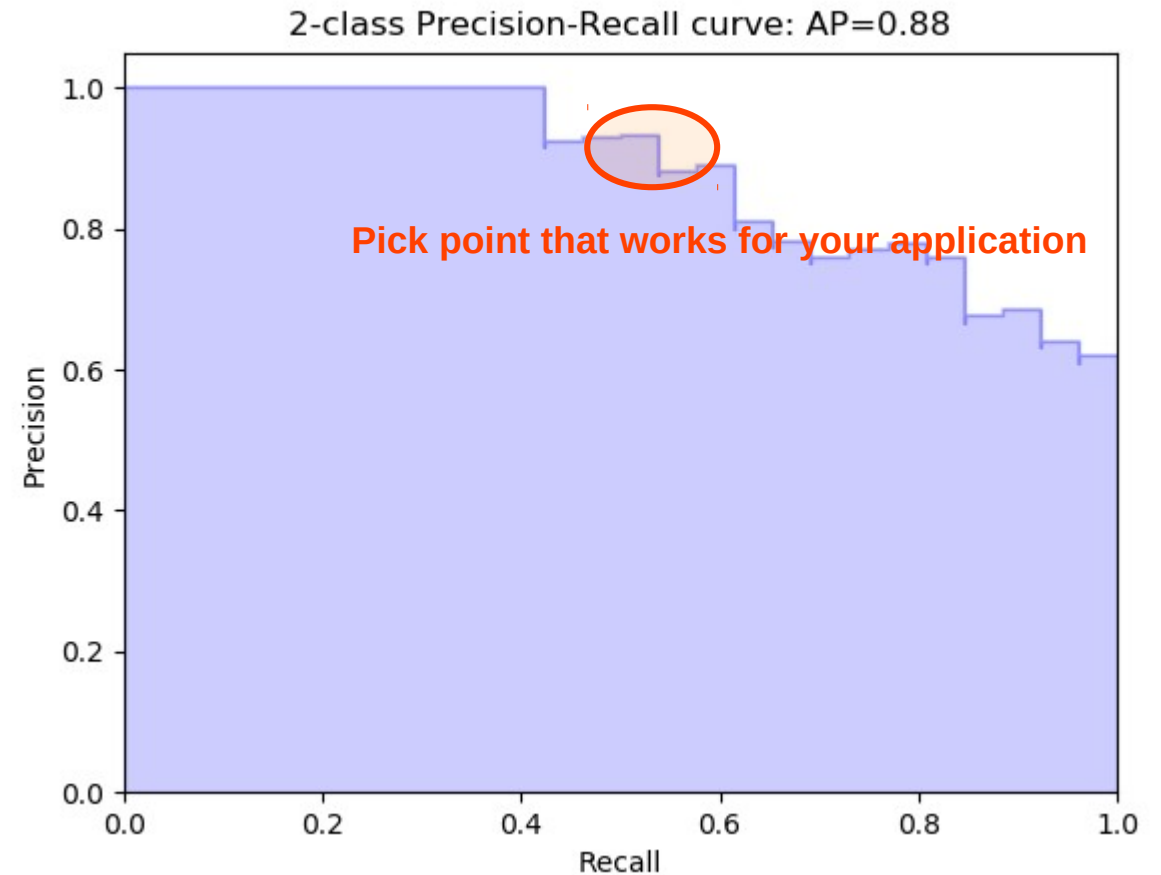
One more note

Often models don't predict 0 or 1

Prediction is a **probability** of being 1

A cutoff has to be chosen such that
probability $>$ CUTOFF means the prediction
and probability \leq CUTOFF means the
prediction is 0

Precision and recall can be computed for
a range of CUTOFFs to give a
precision-recall curve



scikit-learn.org

Decision Tree

Series of if-else statements on features
that distinguish between two classes

Decision Tree

Series of if-else statements on features
that distinguish between two classes


Some we code up all the time!

Decision Tree


Series of if-else statements on features
that distinguish between two classes



Something we code up all the time!

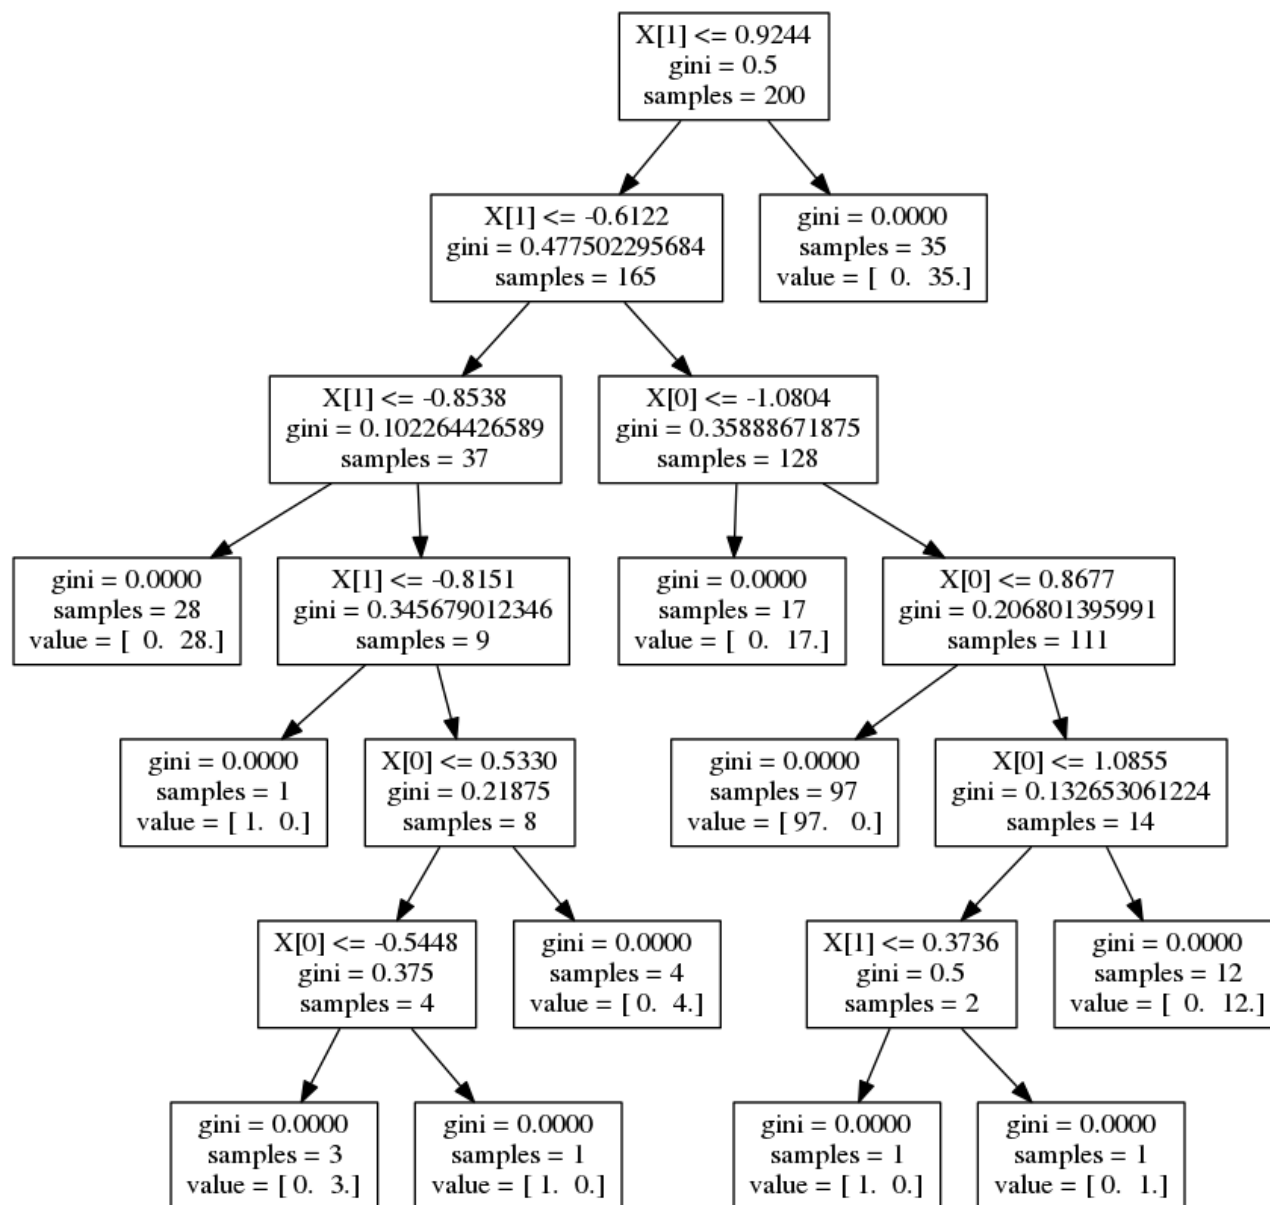
But now the expressions are learned!




Is it possible to separate red and green dots with a sequence of if-else statements on x and y?

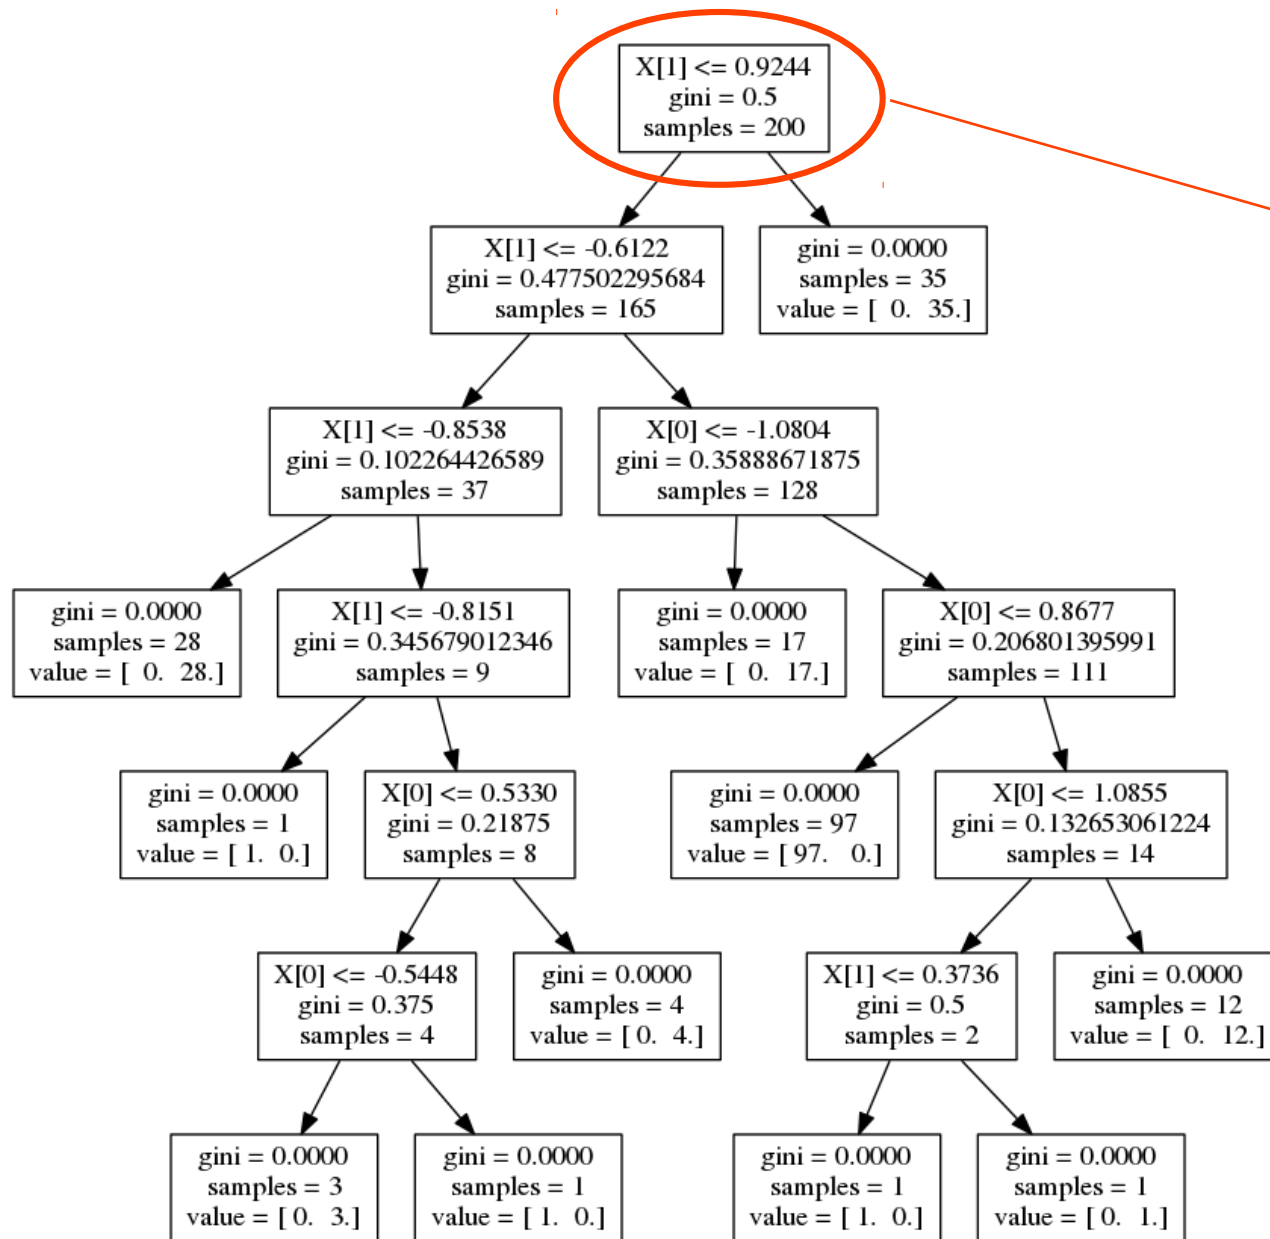








$X[0] = x$
 $X[1] = y$



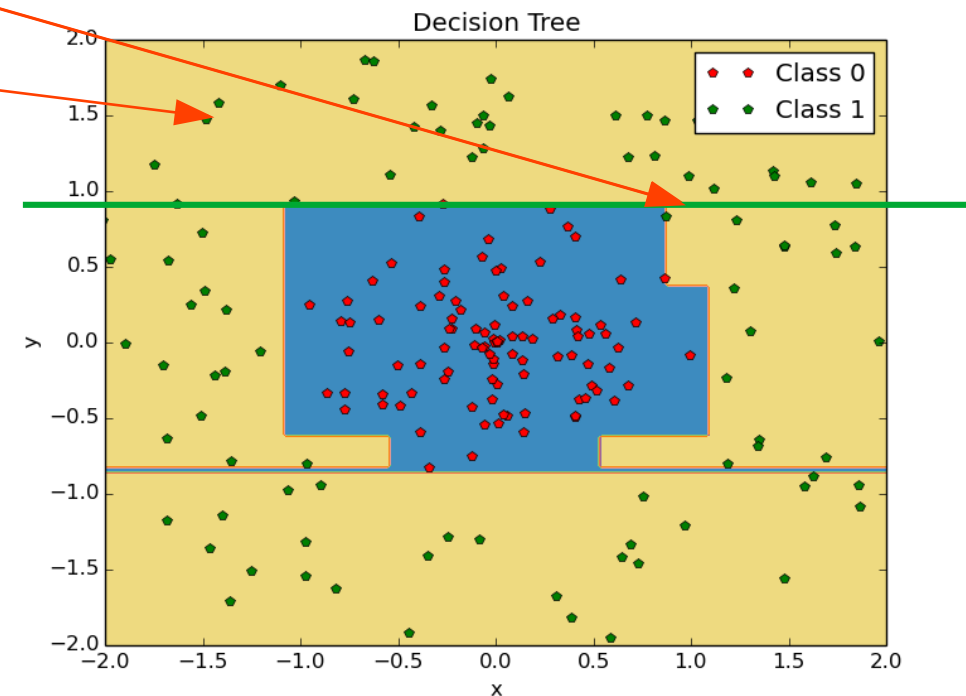


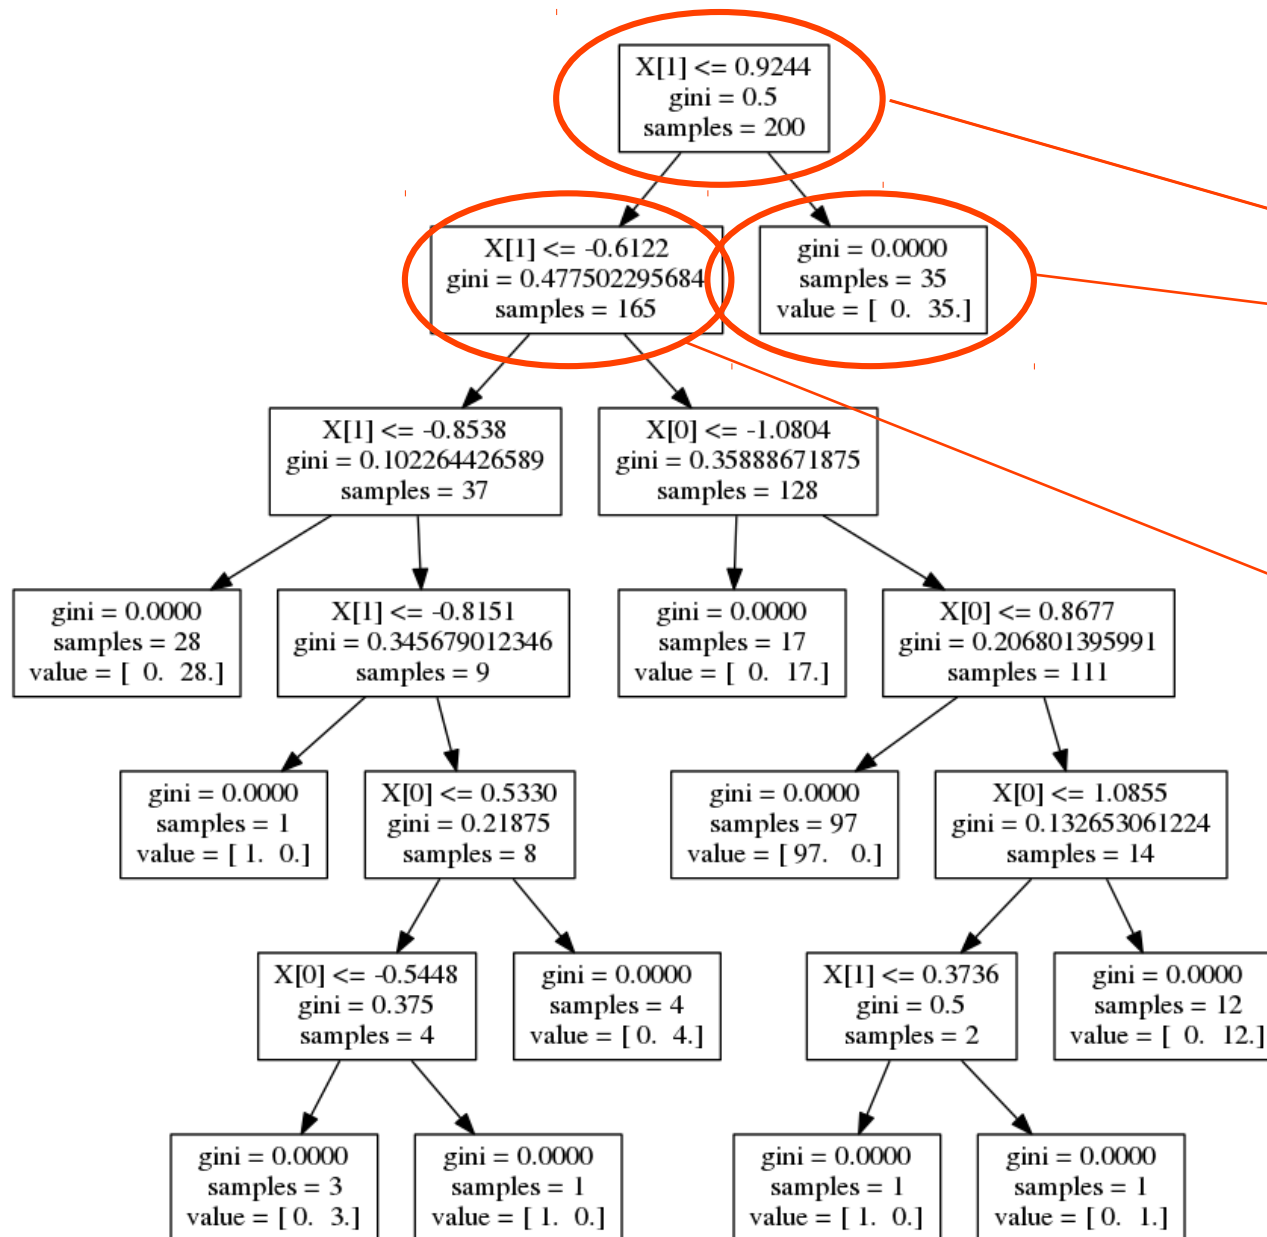
$X[0] = x$
 $X[1] = y$






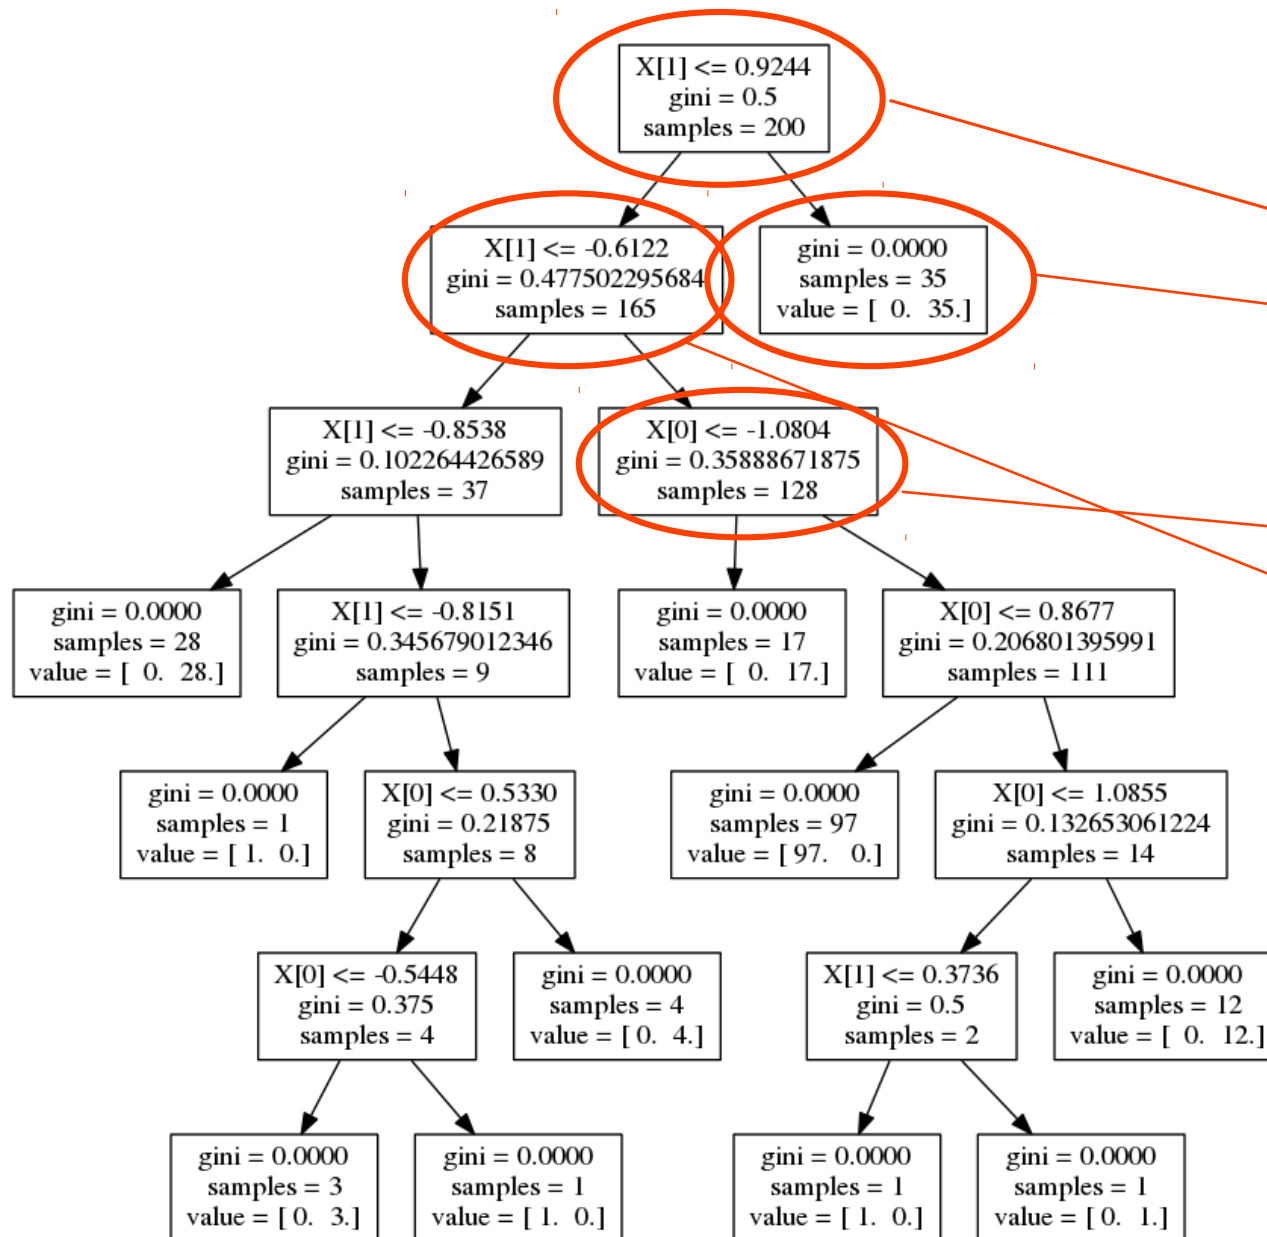
$X[0] = x$
 $X[1] = y$






$X[0] = x$
 $X[1] = y$





$X[0] = x$
 $X[1] = y$



Python Code

```
#Decision Tree Classifier
from sklearn import tree

model = tree.DecisionTreeClassifier()
model.fit(features, labels) #train the model

model.predict(test_features) #make predictions
#See documentation for other functions
```

Decision Trees: Advantages

Minimal pre-processing required - no normalization for example

No assumptions about distribution of data

Simple to interpret - if not too deep

Can discover interactions between features automatically

Run fast

Decision Trees: Disadvantages

Use "horizontal" and "vertical" lines to learn decision boundaries

Result in over-complex trees

Can easily overfit data

Unstable: small changes in data can lead to very different trees

Cannot extrapolate

K-Nearest Neighbors

Given N examples with features x_1, \dots, x_n

Given N results y

New input features:

Find K "closest" examples

Average their results

K-Nearest Neighbors

k (Number of neighbors considered)	Accuracy
1	96.60%
2	95.72%
3	96.63%
4	96.33%
5	96.50%
6	96.42%
7	96.40%
8	96.26%
9	96.16%
10	96.08%


12,600 samples in test set

All neighbors given equal weight – take majority vote

K-Nearest Neighbors

- Labels not uniformly distributed in train data – what if digit “9” occupies 95% of data and the rest only 5% combined. Neighbors biased towards “9”.
- High dimensionality of feature space

Neural Networks




Layered switches that fire (output = 1) if input to switch exceeds threshold, otherwise don't fire (output = 0).

Switches can be combined to form complex logic gates that can “take decisions” based on input

Initially modeled on neurons in brain but anatomically, extremely crude approximation.

Analogy not useful any more.

Support Vector Machines



Pick boundary that creates the largest “gap” between two classes
(Unfortunately, not covered here)

Unsupervised Learning

K-Means Clustering

Define **distance** between two examples: $d(\vec{x}_i, \vec{x}_j) \in \mathbb{R}$



Features for
example i

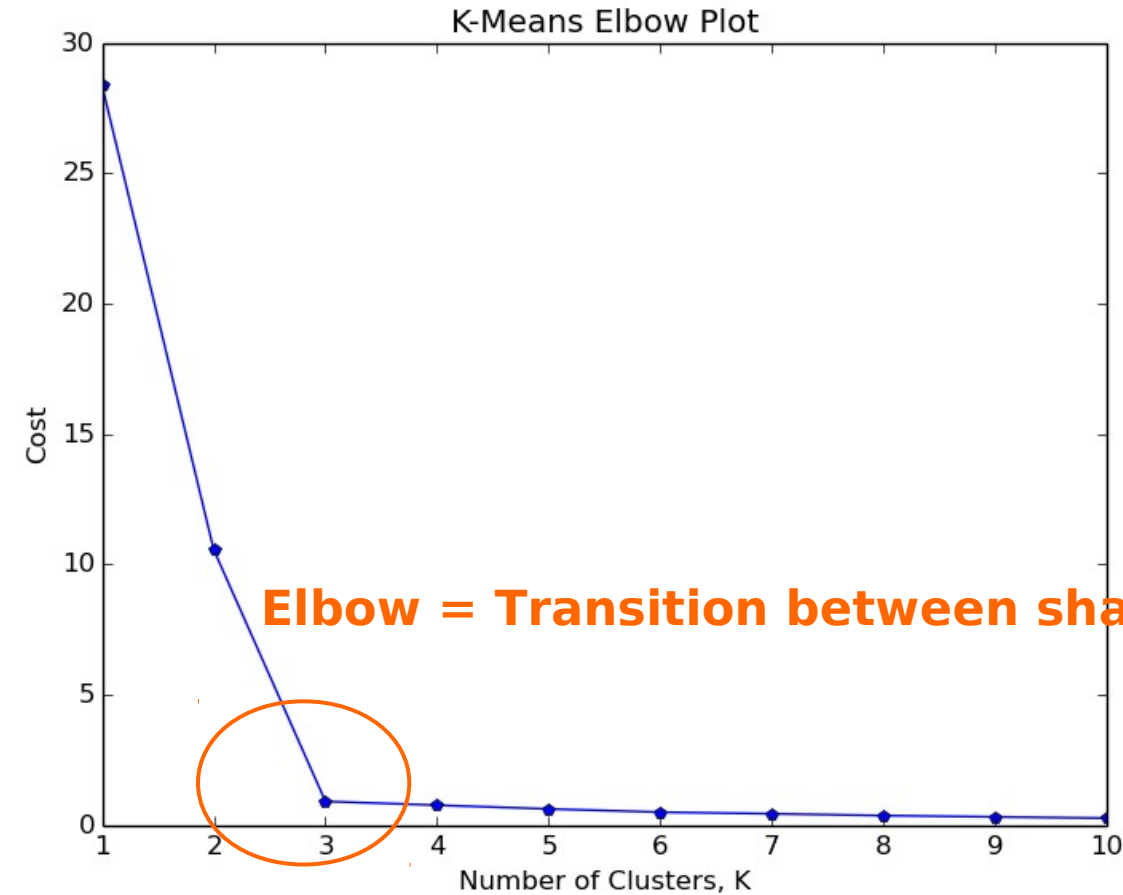
Features for
example j

Analyze/plot each group/cluster separately

75




Choosing k - Elbow Plot



Elbow = Transition between sharp drop and stabilization

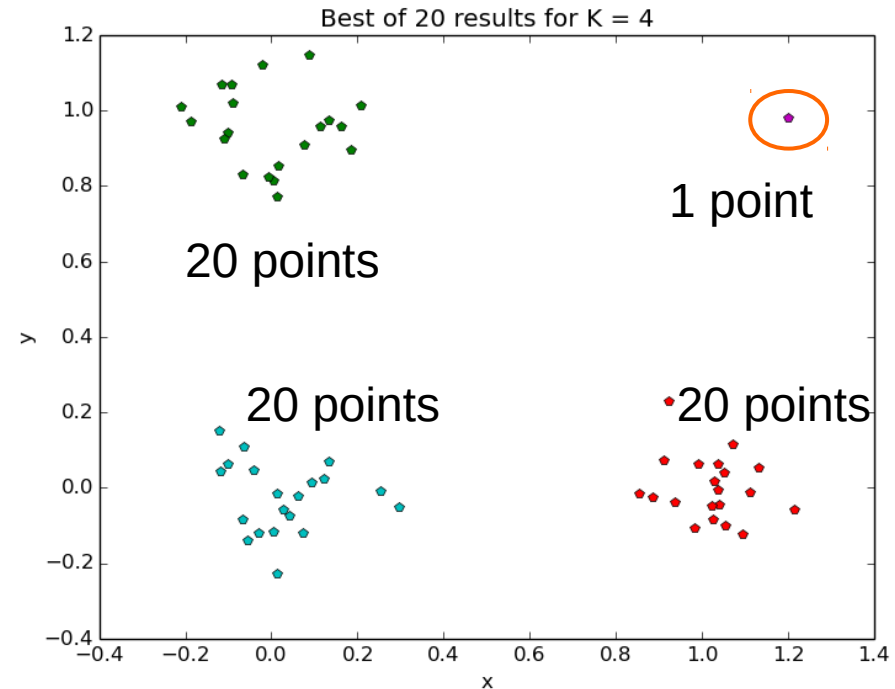
Choose K = 3

Detect Outliers



3 clusters

Cost = 2.05




4 clusters

Cost = 1.19


Best solution over 20 runs

Spectral Clustering



Try K-Means with $K = 2$

Spectral Clustering



Try spectral clustering with $K = 2$

Python Code

```
#Clustering with K-means or Spectral clustering  
from sklearn.cluster import KMeans, SpectralClustering
```

```
model = KMeans(n_clusters = 2)  
model.fit(features) #find clusters  
ids = model.predict(features) #predict cluster IDs
```

k-means

```
model = SpectralClustering(n_clusters = 2)  
model.affinity("nearest_neighbors")  
model.fit(features)  
ids = model.predict(features)
```

Spectral

Warning with Clustering

Good to re-scale each feature to have $\mu = 0$ and $\sigma = 1$

Mean

Standard Deviation

$$x_m \rightarrow x'_m \equiv \frac{x_m - \mu(x_m)}{\sigma(x_m)}$$

Want numbers in similar range for each features

$$\sqrt{\underbrace{(x_1^{(1)} - x_1^{(2)})^2}_{\text{Feature 1}} + \underbrace{(x_2^{(1)} - x_2^{(2)})^2}_{\text{Feature 2}}}$$

Feature 1

Feature 2

If Feature 2 $\approx 10 \times$ Feature 1, distance calculation dominated by Feature 2

$$10^2 + 1000^2 \approx 1000^2$$

Principal Component Analysis

Have n features: x_1, \dots, x_n

Each example is a point in n -dimensional space

Can we visualize the data?


Do we really need all the features?

Example: Plane described by: $4x + 5y - 6z = 1$


2-dimensional surface embedded in 3 dimensions

Need only 2 coordinates not 3

Principal Component Analysis





Random 3-D data set





Projected to 2-D



Principal Component Analysis





Principal Component Analysis





Principal Component Analysis



Principal Component Analysis



Principal Component Analysis



Standard deviation of Gaussian = 2: mixed results

Principal Component Analysis

```
#Principal Component Analysis
from sklearn.decomposition import PCA
model = PCA(n_components = 10)

model.fit(features)
transformed_features = model.transform(features)

model.components_ #new variables
model.explained_variance_ratio_ #% of variance explained by each component
```

Each feature should have mean = 0!

General Modeling Concepts

Choosing the “best” model

You built a model

How do you know if it works?

What does “working” mean?

Typical Scientific Process

Do experiments. Collect data.

Typical Scientific Process

Do experiments. Collect data.

Build model/equations to explain data.

Typical Scientific Process

Do experiments. Collect data.

Build model/equations to explain data.

Make predictions about new experiments.

Typical Scientific Process

Do experiments. Collect data.

Build model/equations to explain data.

Make predictions about new experiments.

Do new experiment. Compare predictions to experimental results.

Typical Scientific Process

Do experiments. Collect data.

Build model/equations to explain data.

Make predictions about new experiments.

Do new experiment. Compare predictions to experimental results.

If results match, model might be correct.

Typical Scientific Process

Do experiments. Collect data.

Build model/equations to explain data.

Make predictions about new experiments.

Do new experiment. Compare predictions to experimental results.

If results don't match, tweak model or build new one.

Typical Scientific Process

Model might work well on already observed data.

After all, we are building the model after looking at this data.

Do experiments. Collect data.

Build model/equations to explain data.

Model might fail terribly when it comes to making predictions, explaining new data

Make predictions about new experiments.

Do new experiment. Compare predictions to experimental results.

If results match, model might be correct.

If results don't match, tweak model or build new one.

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Test data: 30%

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Build model using this data

Test data: 30%

Pretend this doesn't exist

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Build model using this data

**Evaluate performance here
according to some metrics**

Test data: 30%

Pretend this doesn't exist

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Build model using this data

**Evaluate performance here
according to some metrics**

Test data: 30%

Pretend this doesn't exist

**Make predictions here and compute
same metrics**

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Build model using this data

**Evaluate performance here
according to some metrics**

Test data: 30%

Pretend this doesn't exist

**Make predictions here and compute
same metrics**



If metrics have similar values, model works well on “new” test data

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Build model using this data

**Evaluate performance here
according to some metrics**

Test data: 30%

Pretend this doesn't exist

**Make predictions here and compute
same metrics**

If metrics have very different values, model works very well on training data but doesn't capture underlying model well and so works badly on test data - **overfitting**

Model Validation

Have N rows of data with features (inputs) and results (outputs)

Split into two pieces, say 70%-30%
(just a proposal – nothing sacred about these numbers)

Training data: 70%

Build model using this data

**Evaluate performance here
according to some metrics**

Test data: 30%

Pretend this doesn't exist

**Make predictions here and compute
same metrics**

If model performs badly on training data itself - underfitting

Model Validation

Good results on training **AND** test sets

DON'T GUARANTEE

good results on future data

Model Validation

Good results on training **AND** test sets

DON'T GUARANTEE

good results on future data

Assumption 1:

Statistical distributions in train and test set are similar
and will stay so in future data

Assumption 2:

Fundamental rules/processes governing the system
won't change

Cross-Validation

Some models have parameters that can't be decided by training data

Parameters label family of models

Example: Linear Regression with "Regularization"

Penalty term introduced in cost term to prevent overfitting

Free parameter, α

Cross-Validation

Example: Linear Regression with "Regularization"

Penalty term introduced in cost term to prevent overfitting

Free parameter, α

```
model = linear_model.Ridge(alpha = 0.5) # "Ridge" regression with one free parameter  
model = linear_model.Lasso(alpha = 0.5) # "Lasso" regression with one free parameter
```

How do we choose this parameter?

In other words, how do we choose the optimal model from this family of models?

Cross-Validation

Have N rows of data with features and results

Split into three pieces, say 60%-20%-20%

(just a proposal!!!)

Training data: 60%

Test data: 20%

Cross-validation data: 20%

Cross-Validation

Training data: 60%

Train models for different α using this

Cross-validation data: 20%

Apply here

Pick α such that error minimized on cross-validation data

Apply to test data

Model Selection

- Limited amount of data available
- Reuse data (in a controlled fashion)
- K-fold:
 - Split data into k chunk
 - Use k-1 chunks for training, remaining for testing
 - Loop, selecting a different chunk as testing set each time
 - Average errors over all loop iterations

Model Selection – Python Code

```
#Cross-validation  
  
from sklearn import cross_validation  
  
train_features, train_labels, test_features, test_labels = cross_validation.train_test_split(features, labels, test_size = 0.30)  
  
from sklearn.cross_validation import KFold  
kf = KFold(4, n_folds = 4) —————> Split data into 4 chunks. Use 3 for training, 1 for testing, then average.  
for train, test in kf:  
    print train, test  
    train_features, test_features, train_labels, test_labels = features[train], features[test], labels[train], labels[test]  
  
#Do model fitting/predicting here
```

Check out `sklearn.cross_validation`

http://scikit-learn.org/stable/model_selection.html#model-selection

Some Take-away Points

There are many techniques and new ones are always being invented

Guidelines but no strict rules

Be skeptical. Algorithms will always produce **some numbers**

Try techniques on as many datasets as you can!

RH problems

Kaggle problems

data.gov

Some Take-away Points

There are many techniques and new ones are always being invented

Guidelines but no strict rules

Be skeptical. Algorithms will always produce **some numbers**

Try techniques on as many datasets as you can!

RH problems

Kaggle problems

data.gov

Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)



twitter.com/RedHat