

# Linear Models for Regression

# Problem Definition

Given features :  $x_1, x_2, \dots, x_n$

predict real number  $y$


# Simplest Guess

$$y = w_0 + w_1x_1 + \dots + w_nx_n$$

Linear in weights, **NOT** features

For example, maybe  $x_1 = x$  and  $x_n = x^n, n > 1$   
Definitely not linear in features  $x$ , but linear in weights  $w_i$


# Linear Regression: Example



Will see this again. ~1 means good fit

$$y = w_0 + w_1x + w_2x^2 = 0.344 + 0.551x - 0.031x^2$$

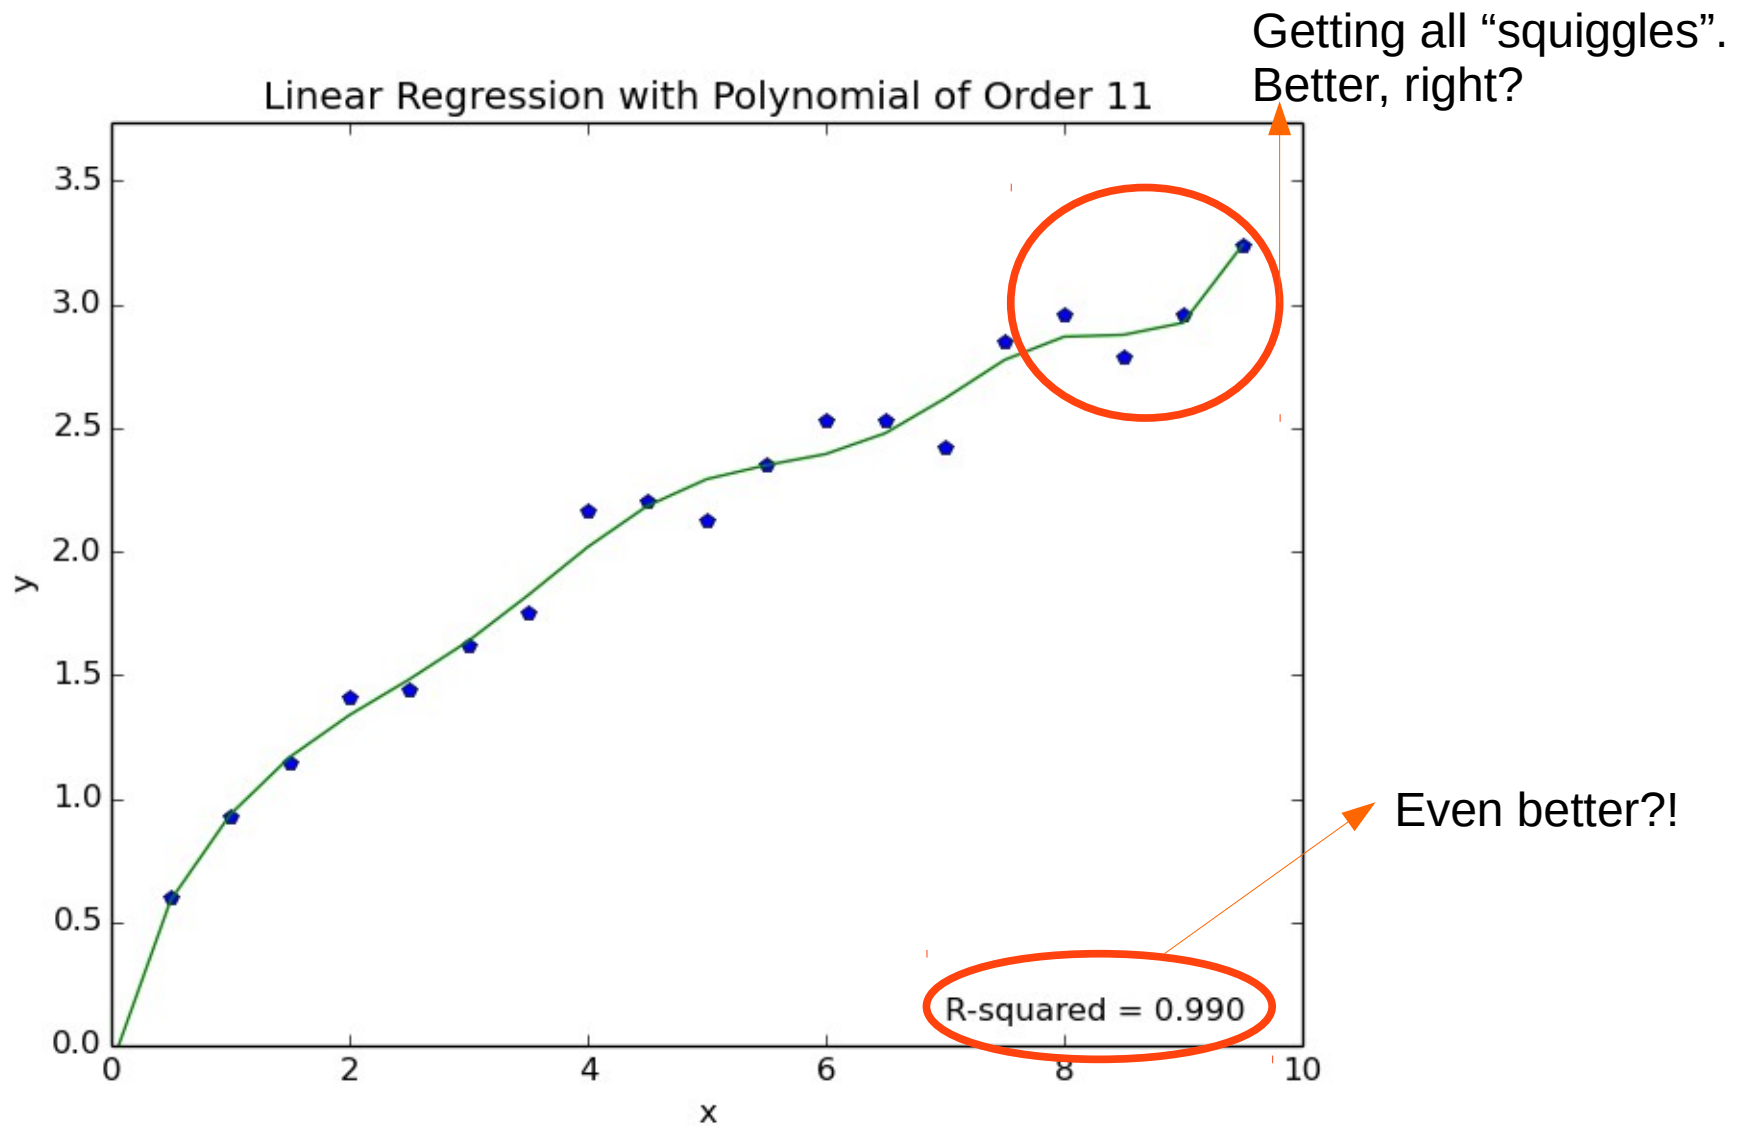
# Linear Regression: Example



$$y = w_0 + w_1x + w_2x^2 + w_3x^2 + w_4x^4$$


$$y = 0.023 + 1.185x - 0.287x^2 + 0.0345x^3 - 0.001x^4$$

# Linear Regression: Example



$$y = w_0 + w_1x + \dots + w_{11}x^{11}$$

# Linear Regression: Example



**Overfitting:** Model too complex and fits too well to training data. Does badly when applied to new data.

# Linear Regression: Training

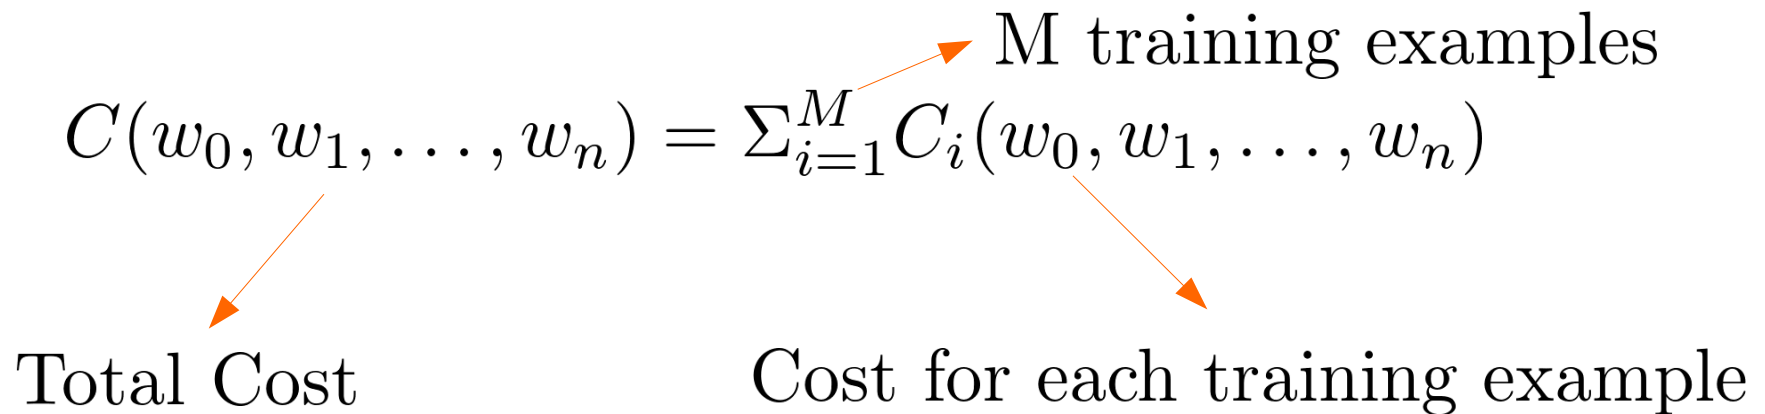
Define cost function to minimize:

$$C(w_0, w_1, \dots, w_n) = \sum_{i=1}^M C_i(w_0, w_1, \dots, w_n)$$

Total Cost

M training examples

Cost for each training example

The diagram shows the equation  $C(w_0, w_1, \dots, w_n) = \sum_{i=1}^M C_i(w_0, w_1, \dots, w_n)$ . An orange arrow points from the left side of the equation to the text "Total Cost". Another orange arrow points from the summation symbol  $\sum_{i=1}^M$  to the text "M training examples". A third orange arrow points from the  $C_i$  term to the text "Cost for each training example".

What should we choose for  $C_i(w_0, w_1, \dots, w_n)$ ?



# Linear Regression: Training

What should we choose for  $C_i(w_0, w_1, \dots, w_n)$ ?

One choice:  $C_i = (y_i - p_i)^2$

$y_i$  = Known value for example i

$p_i$  = Predicted value for example i

$$p_i = w_0 + w_1 x_1^i + \dots + w_n x_n^i$$

 **Example i**

# Linear Regression: Training

Minimize cost function:  $C(w_0, w_1, \dots, w_n)$

$$\text{Solve: } \frac{\partial C}{\partial w_i} = 0, i = 0, \dots, n$$

$$C = \sum_{i=1}^M (y_i - (w_0 + w_1 x_1^i + \dots + w_n x_n^i))^2$$

# Linear Regression: Training

$$\sum_{i=1}^M (y_i - w_0 - w_1 x_1^i - \dots - w_n x_n^i) \cdot 1 = 0$$

$$\sum_{i=1}^M (y_i - w_0 - w_1 x_1^i - \dots - w_n x_n^i) \cdot x_1^i = 0$$

$\vdots$

$$\sum_{i=1}^M (y_i - w_0 - w_1 x_1^i - \dots - w_n x_n^i) \cdot x_n^i = 0$$

# Linear Regression: Training

Define:

$$Y \equiv \sum_{ex=1}^M y_{ex}$$

$$X_{i,j} \equiv \sum_{ex=1}^M x_i^{ex} x_j^{ex}$$

$$Z_i \equiv \sum_{ex=1}^M y_{ex} x_i^{ex}$$

$$X_{i,0} = X_{0,i} \equiv \sum_{ex=1}^M x_i^{ex}$$

$$X_{0,0} = \sum_{i=1}^M 1 = M$$

$$w_0 X_{0,0} + w_1 X_{0,1} + \dots + w_n X_{0,n} = Y$$

$$w_0 X_{1,0} + w_1 X_{1,1} + w_2 X_{1,2} + \dots + w_n X_{1,n} = Z_1$$

$$\vdots$$

$$w_0 X_{n,0} + w_1 X_{n,1} + \dots + w_n X_{n,n} = Z_n$$

# Linear Regression: Training

$$w_0 X_{0,0} + w_1 X_{0,1} + \dots + w_n X_{0,n} = Y$$

$$w_0 X_{1,0} + w_1 X_{1,1} + w_2 X_{1,2} + \dots + w_n X_{1,n} = Z_1$$

$$w_0 X_{n,0} + w_1 X_{n,1} + \dots + w_n X_{n,n} = Z_n$$

**Invert matrix to solve for weights**

$$\begin{pmatrix} X_{0,0} & X_{0,1} & \dots & X_{0,n} \\ X_{1,0} & X_{1,1} & \dots & X_{1,n} \\ \vdots & \vdots & \vdots & \vdots \\ X_{n,0} & X_{n,1} & \dots & X_{n,n} \end{pmatrix} \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_n \end{pmatrix} = \begin{pmatrix} Y \\ Z_1 \\ \vdots \\ Z_n \end{pmatrix}$$



**Depend on data**

# Linear Regression: Assumptions

Data:  $y = w_0 + w_1x_1 + \dots + w_nx_n + \mathcal{N}(0, \sigma^2)$

**Gaussian Noise**

**Linear relationship between  $y$  and  $x_m$ ,  $\forall m$**   
where all other  $x_n$  are fixed

$\mathcal{N}(0, \sigma^2)$  : **Gaussian** noise with **0 mean** and **constant variance**

Might consist of many noise terms

They should be **independent**, **Gaussian** with  
**constant variance**

# Linear Regression: Priors

**No assumptions made about values that weights take**

Minimizing Least Squares Error



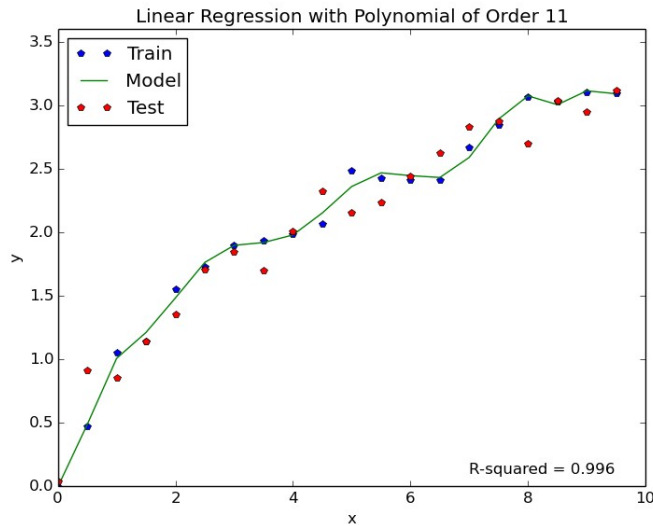
Maximizing Likelihood assuming  
Gaussian predictions

Likelihood,  $LL = \prod_{i=1}^M \frac{1}{2\pi\sigma^2} e^{-\frac{(y_i - w_0 - w_1 x_1 - \dots - w_n x_n)^2}{2\sigma^2}}$

Probability, under a Gaussian distribution, that the model will predict the values observed.

# Linear Regression: Priors

Insert variances



**Recall overfitting example**

Large, alternating weights  
to account for “squigglyness”

$w_0$	-0.00908524896296
$w_1$	-1.68602215
$w_2$	11.8480971
$w_3$	-19.3753732
$w_4$	16.1142291
$w_5$	-7.88218507
$w_6$	2.41735237
$w_7$	-0.477580238
$w_8$	0.0606973836
$w_9$	-0.00479240240
$w_{10}$	0.00023823655
$w_{11}$	-0.00000411844556

Please excuse significant figures!



# Linear Regression: L2 Regularization (Ridge)

Assumption about weights: follow Gaussian distribution.

$$\frac{1}{2\pi\tau^2} e^{-\frac{w_i^2}{2\tau^2}}, i = 1, \dots, n$$

Penalty on weights away from 0

$$LL = \prod_{i=1}^M \frac{1}{2\pi\sigma^2} e^{-\frac{(y_i - w_0 - w_1 x_1 - \dots - w_n x_n)^2}{2\sigma^2}} \prod_{i=1}^n \frac{1}{2\pi\tau^2} e^{-\frac{w_i^2}{2\tau^2}}$$

Maximizing Likelihood



Minimizing new cost function

$$C = \sum_{i=1}^M (y_i - (w_0 + w_1 x_1^i + \dots + w_n x_n^i))^2 + \frac{\sigma^2}{\tau^2} \sum_{i=1}^n w_i^2$$

Prevents overfitting

# Linear Regression: L1 Regularization (Lasso)

Assumption about weights: follow Laplace distribution.

$$\frac{\lambda}{2} e^{-\lambda |w_i|}, i = 1, \dots, n$$

Penalty on weights away from 0

$$LL = \prod_{i=1}^M \frac{1}{2\pi\sigma^2} e^{-\frac{(y_i - w_0 - w_1 x_1 - \dots - w_n x_n)^2}{2\sigma^2}} \prod_{i=1}^n \frac{\lambda}{2} e^{-\lambda |w_i|}$$

Maximizing Likelihood



Minimizing new cost function

$$C = \sum_{i=1}^M (y_i - (w_0 + w_1 x_1^i + \dots + w_n x_n^i))^2 + 2\lambda\sigma^2 \sum_{i=1}^n |w_i|$$

Prevents overfitting

# Linear Regression: Python Code

```
from sklearn import linear_model

model = linear_model.LinearRegression() #create instance of linear regression model
model.fit(x_values, y_values) #train model

#DONE!!!

model.predict(new_x_values) #predict result on new inputs

model.coef_ #see weights

model.score(x_values, y_values) #we'll see this soon. Compute R2
```

# Linear Regression: Python Code

```
model = linear_model.Ridge(alpha = 0.5) #"Ridge" regression with one free parameter  
model = linear_model.Lasso(alpha = 0.5) #"Lasso" regression with one free parameter
```

Use these instead of linear model used previously  
Prevent overfitting (coming up soon)

Also, it's very enlightening to look at source code if interested

# Linear Regression: Scaling

Generally, features  $x_m$  can have different ranges they take values in

Good to normalized each feature to have mean = 0 and standard deviation = 1

$$x_m \rightarrow \frac{x_m - \mu(x_m)}{\sigma(x_m)}$$

Mean across all examples

Standard dev. across all examples

Mean = 0 helps with convergence of algorithms too

according to influence on output

# Linear Regression: Performance

$$\text{Residual} \equiv \sum_{i=1}^N (y_i - f(x_i))^2$$

Result of  
example i

Model prediction for  
example i

$$\text{Variance} \equiv \sum_{i=1}^N (y_i - \bar{y})^2$$

$$\text{Unexplained Variance Ratio} = \frac{\text{Residual}}{\text{Variance}}$$


$$R^2 \equiv 1 - \frac{\text{Residual}}{\text{Variance}}$$

Close to 1 = Good(\*)

# Linear Regression: Performance

$$R^2 \equiv 1 - \frac{\text{Residual}}{\text{Variance}}$$

Close to 1 but might be overfitting



Compare  $R^2$  between train and test sets!!!!

# Linear Regression: Performance

Simpler model is better

$R^2_{\text{Adjusted}}$  penalizes model with many parameters

$$R^2 \equiv 1 - \frac{\text{Residual Variance}}{\text{Variance}}$$



$$R^2_{\text{Adjusted}} \equiv 1 - \frac{\frac{\text{Residual}}{df_e}}{\frac{\text{Variance}}{df_t}}$$

$$df_t \equiv n - 1 \longrightarrow \text{"Degrees of freedom for variance"}$$

$$df_e \equiv n - p - 1 \longrightarrow \text{"Degree of freedom for model"}$$

$n$  = Number of examples

$p$  = number of weights (excluding constant term)



# Linear Regression: Performance

Look at distribution of residuals:  $y_i - f(x_i)$

Should be Gaussian with constant variance

# Linear Models for Classification: Logistic “Regression”

# Problem Definition

Given features :  $x_1, x_2, \dots, x_n$

predict class or label,  $y = 0$  or  $1$

# First Stab

Group together all distinct values of tuples  $(x_1, x_2, \dots, x_n)$

For each distinct group, count number of entries with  $y = 0$  and  $y = 1$

Now have  $\text{Prob}[y = 0 | (x_1, x_2, \dots, x_n)]$

$$\text{Prob}[y = 0 | (x_1, x_2, \dots, x_n)] > P \rightarrow \text{label} = 0$$

$$\text{Prob}[y = 0 | (x_1, x_2, \dots, x_n)] \leq P \rightarrow \text{label} = 1$$



P tuning parameter, e.g., 0.5

# Problems with First Stab

$x_1, x_2, \dots, x_n$  continuous - can't count!

Solution: bucket/bin continuous variables in intervals

Receive new features:  $(x_1, x_2, \dots, x_n)$

NOT in training set

Have no data on  $Prob[y = 0 | (x_1, x_2, \dots, x_n)]!!!$

Solution: Find closest  $(x_1, x_2, \dots, x_n)$  that exists in training set  
and use it as a proxy

# Linear Model

Can we map:

$$(x_1, x_2, \dots, x_n)$$



$$Prob[y|(x_1, x_2, \dots, x_n)]$$

using some other method?

# Linear Model


How about:

$$(x_1, x_2, \dots, x_n)$$



$$w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

Linear combination:



$$[0, 1]$$

# Logistic/Sigmoid Function

Want to map real number to  $[0,1]$

One possible solution:  $f(x) = \frac{1}{1 + e^{-x}}$

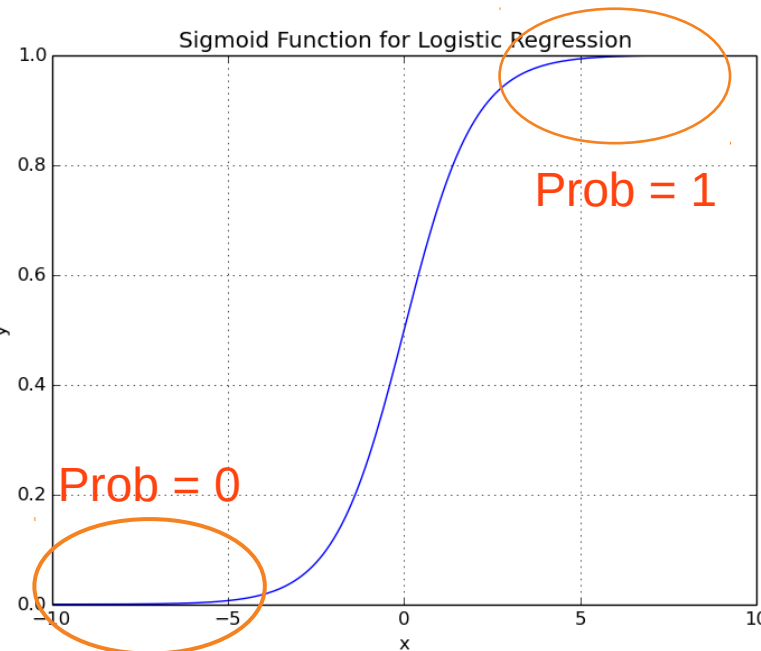
$$x \rightarrow \infty, f(x) \rightarrow \frac{1}{1 + 0} = 1$$

$$x \rightarrow -\infty, f(x) \rightarrow \frac{1}{1 + \infty} = 0$$

$$x = 0, f(x) = 0.5$$



# Logistic/Sigmoid Function



$$x \rightarrow -\infty, f(x) \rightarrow \frac{1}{1 + \infty} = 0$$

$$x \rightarrow \infty, f(x) \rightarrow \frac{1}{1 + 0} = 1$$

Probability

Want to map real number to  $[0,1]$

One possible solution:  $f(x) = \frac{1}{1 + e^{-x}}$

$x = 0, f(x) = 0.5 \longrightarrow$  "Boundary" or "Threshold"

# Logistic Regression

$$(x_1, x_2, \dots, x_n)$$



$$w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n$$

Linear combination:



$$f(w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n) = \frac{1}{1 + e^{-(w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n)}}$$

< 0.5

= 0.5

> 0.5



Label 0




Boundary separating classes:



Label 1

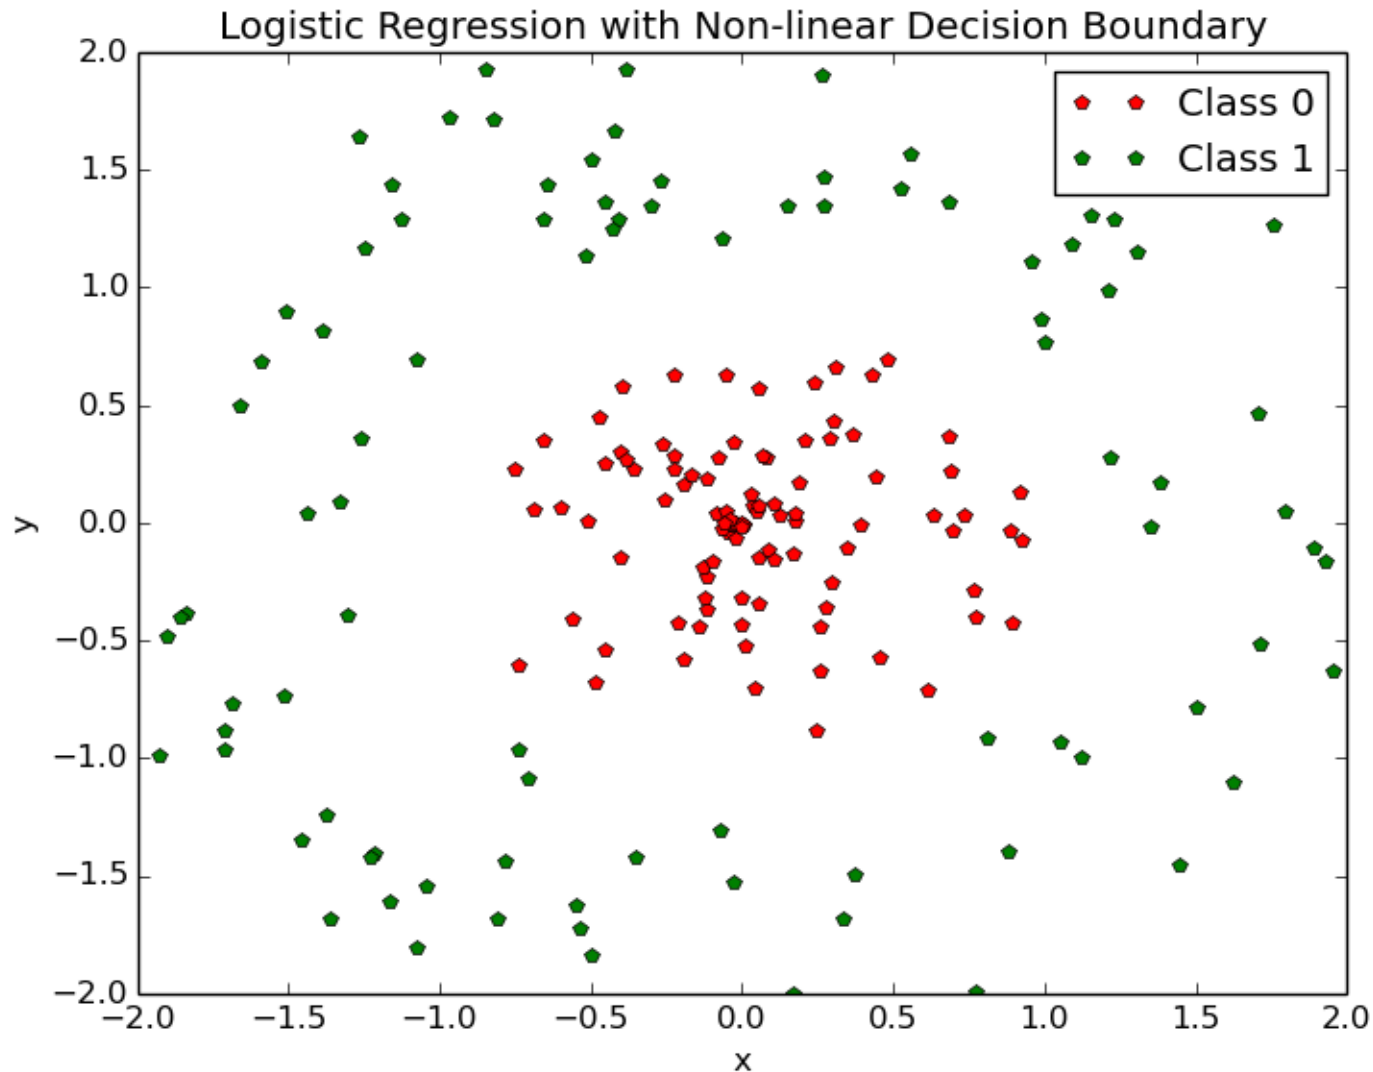
$$w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n = 0$$

# Logistic Regression Example 1




Decision Boundary :  $-4.86 + 1.05x + 1.00y = 0$

# Logistic Regression Example 2



**Accuracy = 64%**

# Logistic Regression Example 2




Bad!!!

Decision Boundary:  $w_0 + w_1x + w_2y = 0 \implies$  Straight Line

# Logistic Regression: Example 2

Features are:  $(x^2, y^2)$  NOT  $(x, y)$




**Accuracy = 100%**

Decision Boundary:  $w_0 + w_1x^2 + w_2y^2 = 0 \implies$  Ellipse (in this case at least)

# Logistic Regression: Example 2

Features are:  $(\log |x|, \log |y|)$  NOT  $(x, y)$



Accuracy = 86%

Decision Boundary:  $a \log |x| + b \log |y| + c = 0 \implies y = \pm \frac{e^{-\frac{c}{b}}}{|x|^{\frac{a}{b}}}$

# Logistic Regression: Training

Define cost function to minimize:

$$C(w_0, w_1, \dots, w_n) = \sum_{i=1}^M C_i(w_0, w_1, \dots, w_n)$$

Total Cost

M training examples

Cost for each training example

What should we choose for  $C_i(w_0, w_1, \dots, w_n)$ ?



# Logistic Regression: Training

What should we choose for  $C_i(w_0, w_1, \dots, w_n)$ ?

One choice:  $C_i = (y_i - p_i)^2$

$y_i$  = Known class of example  $i \in \{0, 1\}$

$p_i$  = Predicted class of example  $i \in \{0, 1\}$

Could also try:

$p_i$  = Predicted probability of example  $i = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + \dots + w_n x_n)}}$

# Logistic Regression: Training

One choice:  $C_i = (y_i - p_i)^2$

$C_i$	$y_i = 0$	$y_i = 1$
$p_i = 0$	0	1
$p_i = 1$	1	0

Bounded cost  
even when result  
wrong

Unlike linear regression, in logistic regression, the result is binary: right or wrong.

Want to impose high cost for wrong and no cost for right.

Let's try something else.

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

where:

$y_i$  = Known class of example  $i \in \{0, 1\}$

$p_i$  = Predicted probability of example  $i = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + \dots + w_n x_n)}}$

Why this  $C_i$ ?!!!

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 0$$

$$p_i = 0$$

$$C_i = 0?$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 0$$

$$p_i = 0$$

$$C_i = 0?$$

$$C_i = -0 \log 0 - (1 - 0) \log(1 - 0)$$

$$= 0 - \log 1 = 0$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 1$$

$$p_i = 1$$

$$C_i = 0?$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 1$$

$$p_i = 1$$

$$C_i = 0?$$

$$C_i = -1 \log 1 - (1 - 1) \log(1 - 1)$$

$$= -\log 1 - 0 \log 0 = 0$$



# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 1$$

$$p_i = 0$$

$$C_i = \text{large?}$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 1$$

$$p_i = 0$$

$$C_i = \text{large?}$$

$$C_i = -1 \log 0 - (1 - 1) \log (1 - 0)$$

$$\log 0 = -\infty$$

$$C_i = \infty$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 0$$

$$p_i = 1$$

$$C_i = \text{large?}$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

$$y_i = 0$$

$$p_i = 1$$

$$C_i = \text{large?}$$

$$C_i = -0 \log 1 - (1 - 0) \log (1 - 1)$$

$$\log 0 = -\infty$$

$$C_i = \infty$$

# Logistic Regression: Training

Different Choice:  $C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$

Takes values  
between 0 and 1.

Only extreme cases shown.

$C_i$	$y_i = 0$	$y_i = 1$
$p_i = 0$	0	$\infty$
$p_i = 1$	$\infty$	0

**Unbounded** cost  
when result  
wrong

**More appropriate for binary classification**

**Cost function measures “entropy”**

**Also, equivalent to negative log likelihood**

# Logistic Regression: Training

Minimize:

$$C(w_0, w_1, \dots, w_n) = \sum_{i=1}^M C_i(w_0, w_1, \dots, w_n)$$


$$C_i = -y_i \log p_i - (1 - y_i) \log (1 - p_i)$$

$$p_i = \text{Predicted probability of example } i = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + \dots + w_n x_n)}}$$

with respect to:

$$w_0, w_1, \dots, w_n$$

# Logistic Regression: Visualization



**This kind of structure will occur again when we look at neural networks**

# Logistic Regression: Python Code

```
from sklearn import linear_model

model = linear_model.LogisticRegression()
model.fit(features, labels) #train the model

model.predict(test_features) #make predictions (0/1)
model.predict_proba(test_features) #predict probabilities
```



# Logistic Regression: Performance

Have data with binary labels, built a logistic regression model.

How good is the model? How to evaluate its performance?

# Logistic Regression: Performance

Good

Bad

$C_i$	$pred = 0$	$pred = 1$
$label = 0$	$0 \rightarrow 0$	$0 \rightarrow 1$
$label = 1$	$1 \rightarrow 0$	$1 \rightarrow 1$

=

$C_i$	$pred = 0$	$pred = 1$
$label = 0$	True Negatives	False Positives
$label = 1$	False Negatives	True Positives

# Logistic Regression: Performance

**Some metrics to measure on the test set**

$$\frac{(0 \rightarrow 0) + (1 \rightarrow 1)}{(0 \rightarrow 0) + (1 \rightarrow 1) + (0 \rightarrow 1) + (1 \rightarrow 0)}$$

=

$$\frac{\text{examples predicted correctly}}{\text{total examples}}$$

known as **Accuracy**

# Logistic Regression: Performance

$$\text{Accuracy} = \frac{(0 \rightarrow 0) + (1 \rightarrow 1)}{(0 \rightarrow 0) + (1 \rightarrow 1) + (0 \rightarrow 1) + (1 \rightarrow 0)}$$

What if 98% of examples are labeled 0 and 2% labeled 1?

Build "dumb" model that predicts 0 for every examples

Accuracy = 98% (Get everything labeled 0 correct)

Bad measure if population of labeled classes very uneven

# Logistic Regression: Performance

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)}$$

Ideally, Recall = 1

In previous example, Recall = 0

# Logistic Regression: Performance

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)}$$

$$\text{Specificity} = \frac{(0 \rightarrow 0)}{(0 \rightarrow 0) + (0 \rightarrow 1)}$$

Class labels arbitrary

Symmetry under:  $0 \leftrightarrow 1$

Sensitivity/Recall  $\leftrightarrow$  Specificity

**Ideally, both are 1**

# Logistic Regression: Performance

$$\text{type I error} = \frac{(0 \rightarrow 1)}{(0 \rightarrow 1) + (0 \rightarrow 0)}$$


$$0 \leftrightarrow 1 \updownarrow$$

$$\text{type II error} = \frac{(1 \rightarrow 0)}{(1 \rightarrow 0) + (1 \rightarrow 1)}$$

**Ideally, both are 0**

# Logistic Regression: Performance

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)}$$

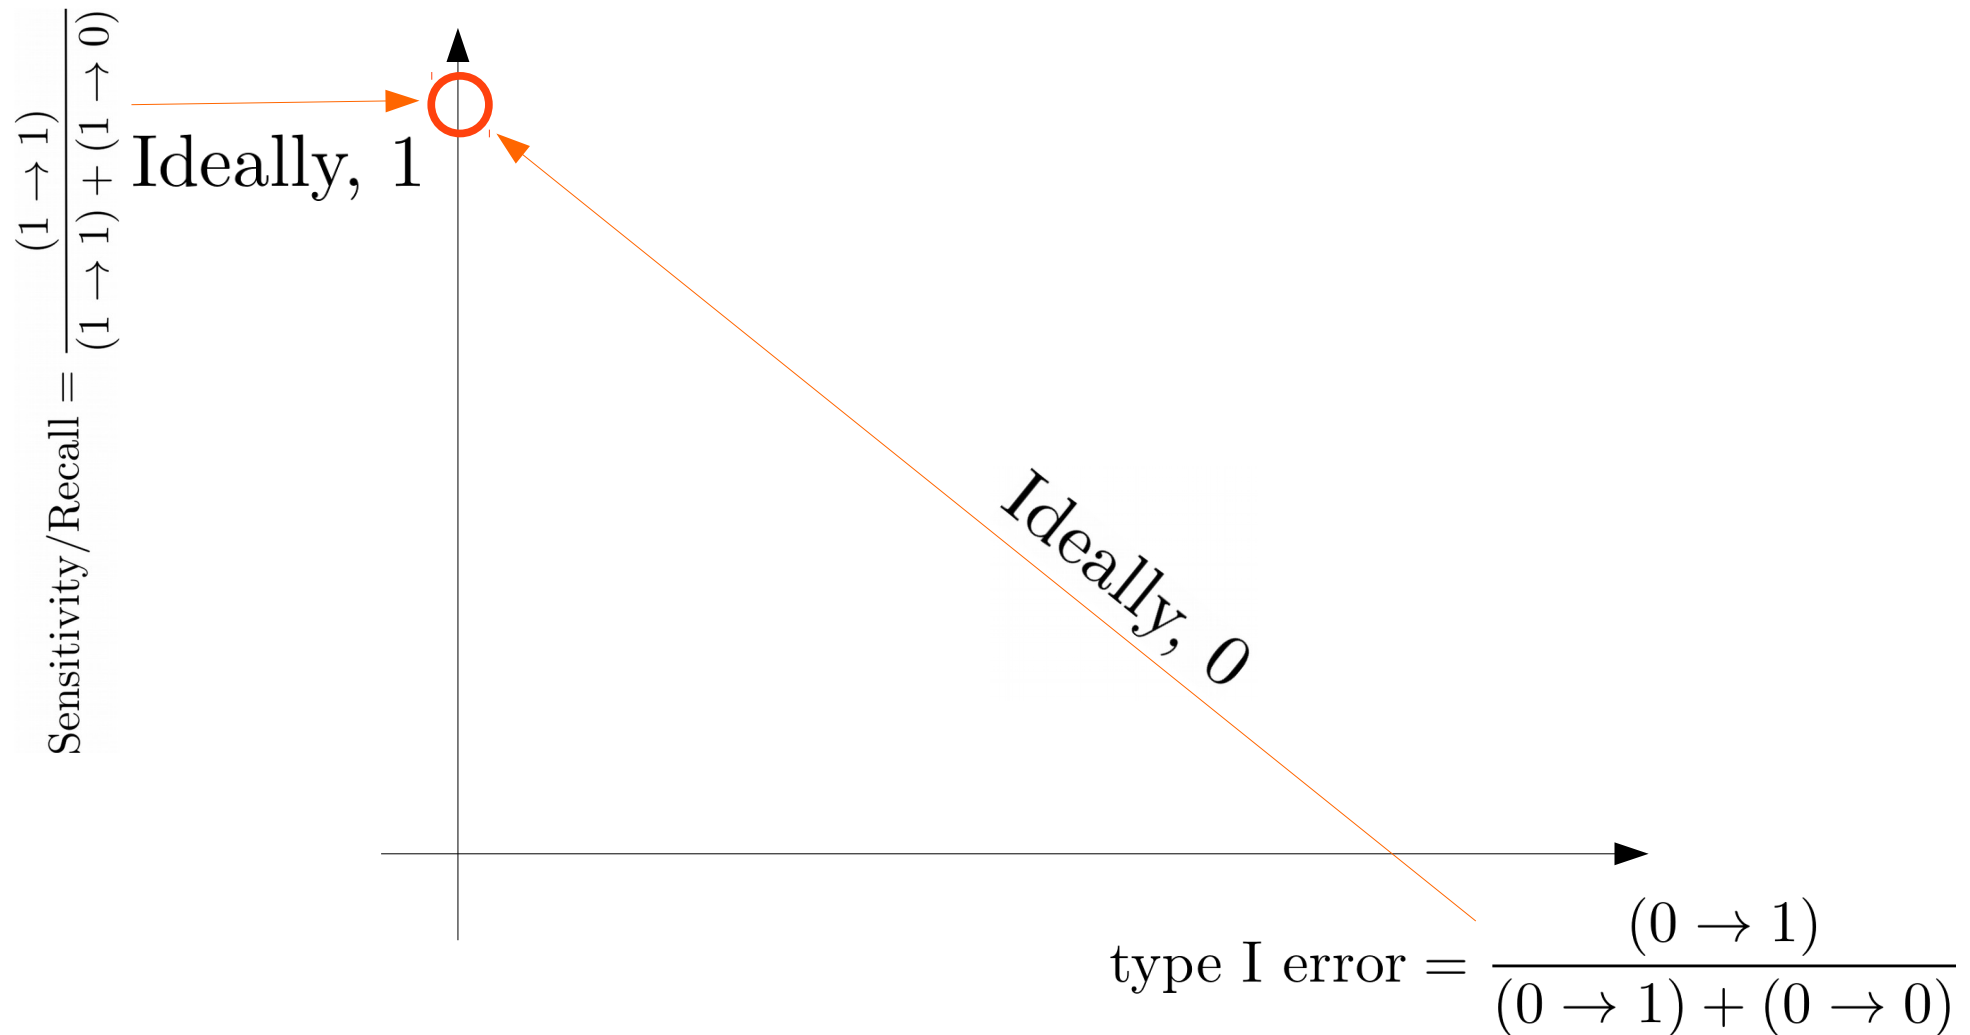


$$\text{type I error} = \frac{(0 \rightarrow 1)}{(0 \rightarrow 1) + (0 \rightarrow 0)}$$

**Receiver Operating Characteristic (ROC) Curve**




# Logistic Regression: Performance




**Receiver Operating Characteristic (ROC) Curve**

# Logistic Regression: Performance



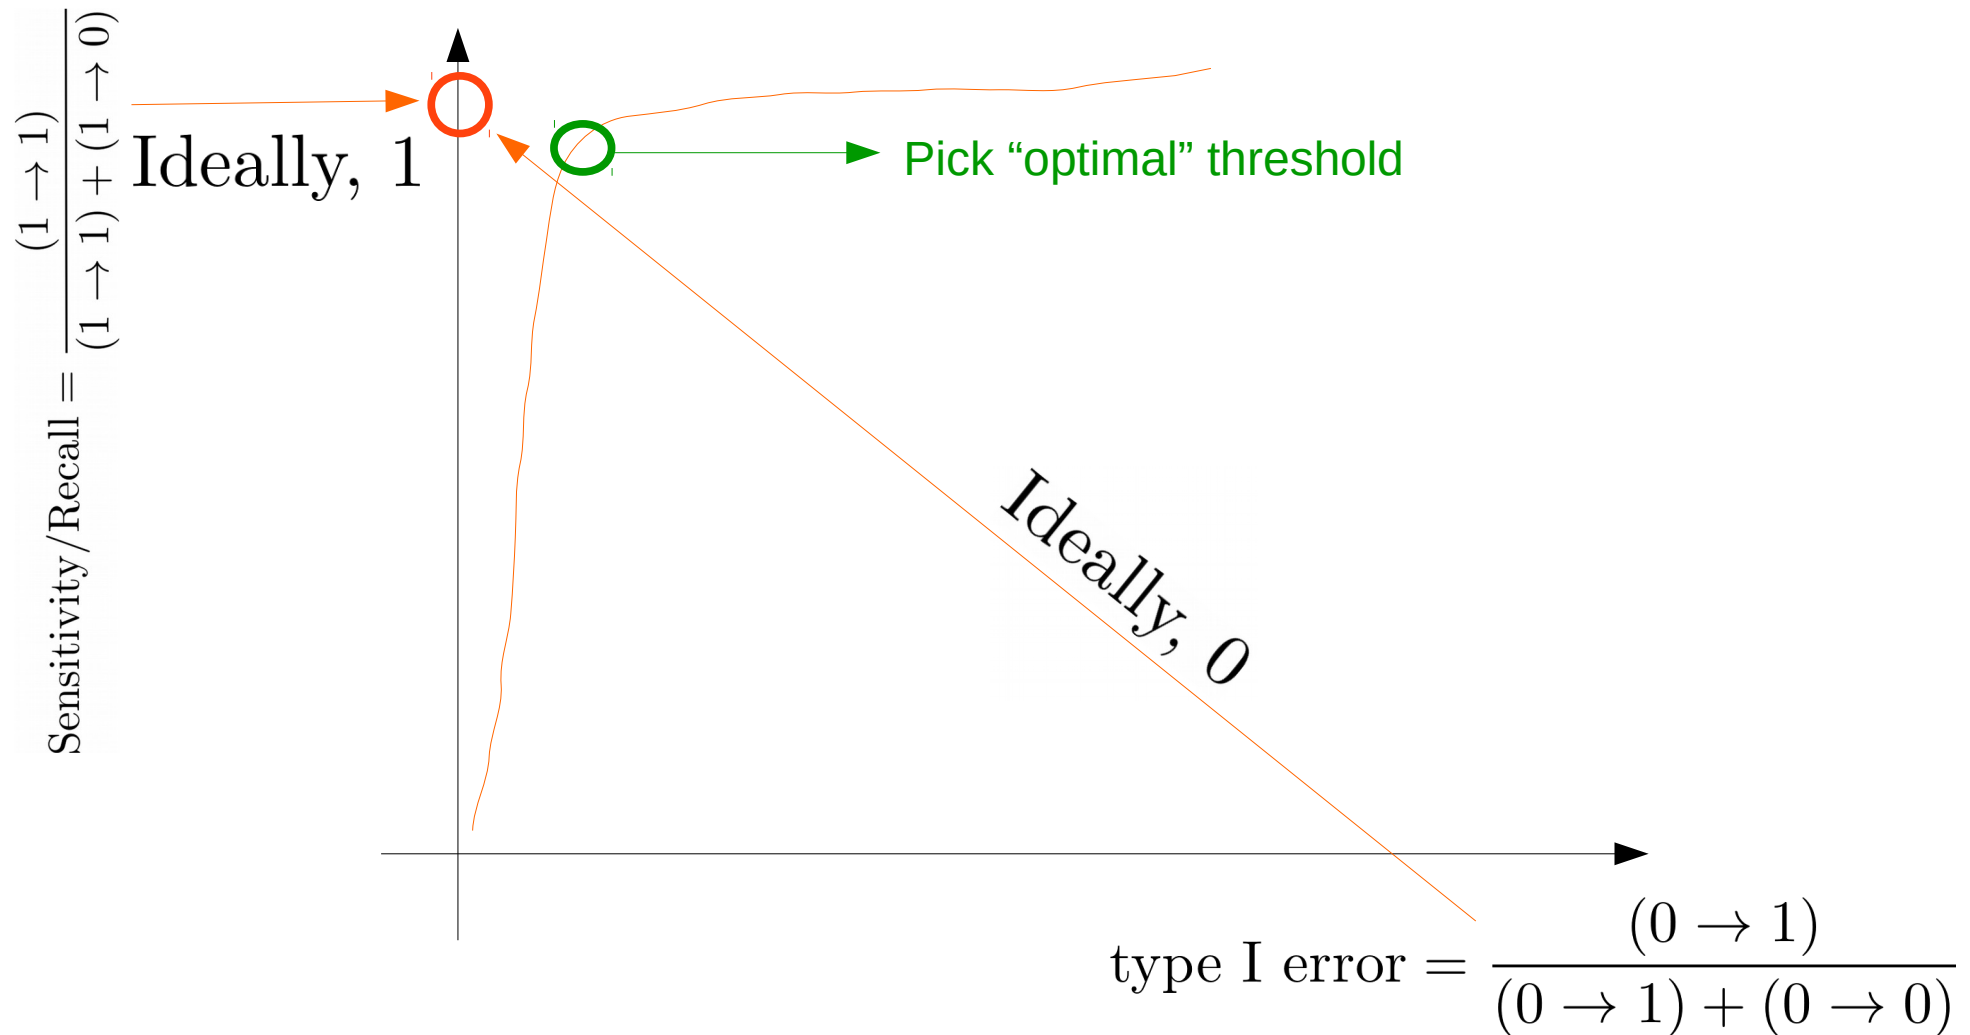
**Receiver Operating Characteristic (ROC) Curve**

# Logistic Regression: Performance




**Receiver Operating Characteristic (ROC) Curve**

# Logistic Regression: Performance



**Receiver Operating Characteristic (ROC) Curve**

# Logistic Regression: Example 3



Hand-written digits dataset: 28 x 28 pixels

Let's try logistic regression to identify the digit 3 (label = 1)  
from other digits (label = 0)

# Logistic Regression: Example 3

Test Set: # of examples with label 1 = 1,301

Test Set: # of examples with label 0 = 11,299

Assign every test example to label = 0:

$$\text{Accuracy} = \frac{11,299}{1,301 + 11,299} = 89.7\%$$

Logistic Regression:

$$\text{Accuracy} = \frac{12,229}{12,600} = 97.1\%$$

# Logistic Regression: Example 3

$$\text{Sensitivity/Recall} = \frac{(1 \rightarrow 1)}{(1 \rightarrow 1) + (1 \rightarrow 0)}$$

Assign every test example to label = 0:

$$\text{Sensitivity/Recall} = \frac{0}{0 + 1301} = 0$$

Logistic Regression:

$$\text{Sensitivity/Recall} = \frac{1,077}{1,077 + 224} = 82.7\%$$

# Logistic Regression: Example 3

$$\text{Specificity} = \frac{(0 \rightarrow 0)}{(0 \rightarrow 0) + (0 \rightarrow 1)}$$

Assign every test example to label = 0:

$$\text{Specificity} = \frac{11,299}{11,299 + 0} = 100\%$$

Logistic Regression:

$$\text{Specificity} = \frac{11,152}{11,152 + 147} = 98.7\%$$



# Logistic Regression: Example 3

$$\text{type I error} = \frac{(0 \rightarrow 1)}{(0 \rightarrow 1) + (0 \rightarrow 0)}$$

Assign every test example to label = 0:

$$\text{type I error} = \frac{0}{0 + 11,299} = 0\%$$

Logistic Regression:

$$\text{type I error} = \frac{147}{147 + 11,152} = 1.3\%$$

# Logistic Regression: Example 3

$$\text{type II error} = \frac{(1 \rightarrow 0)}{(1 \rightarrow 0) + (1 \rightarrow 1)}$$

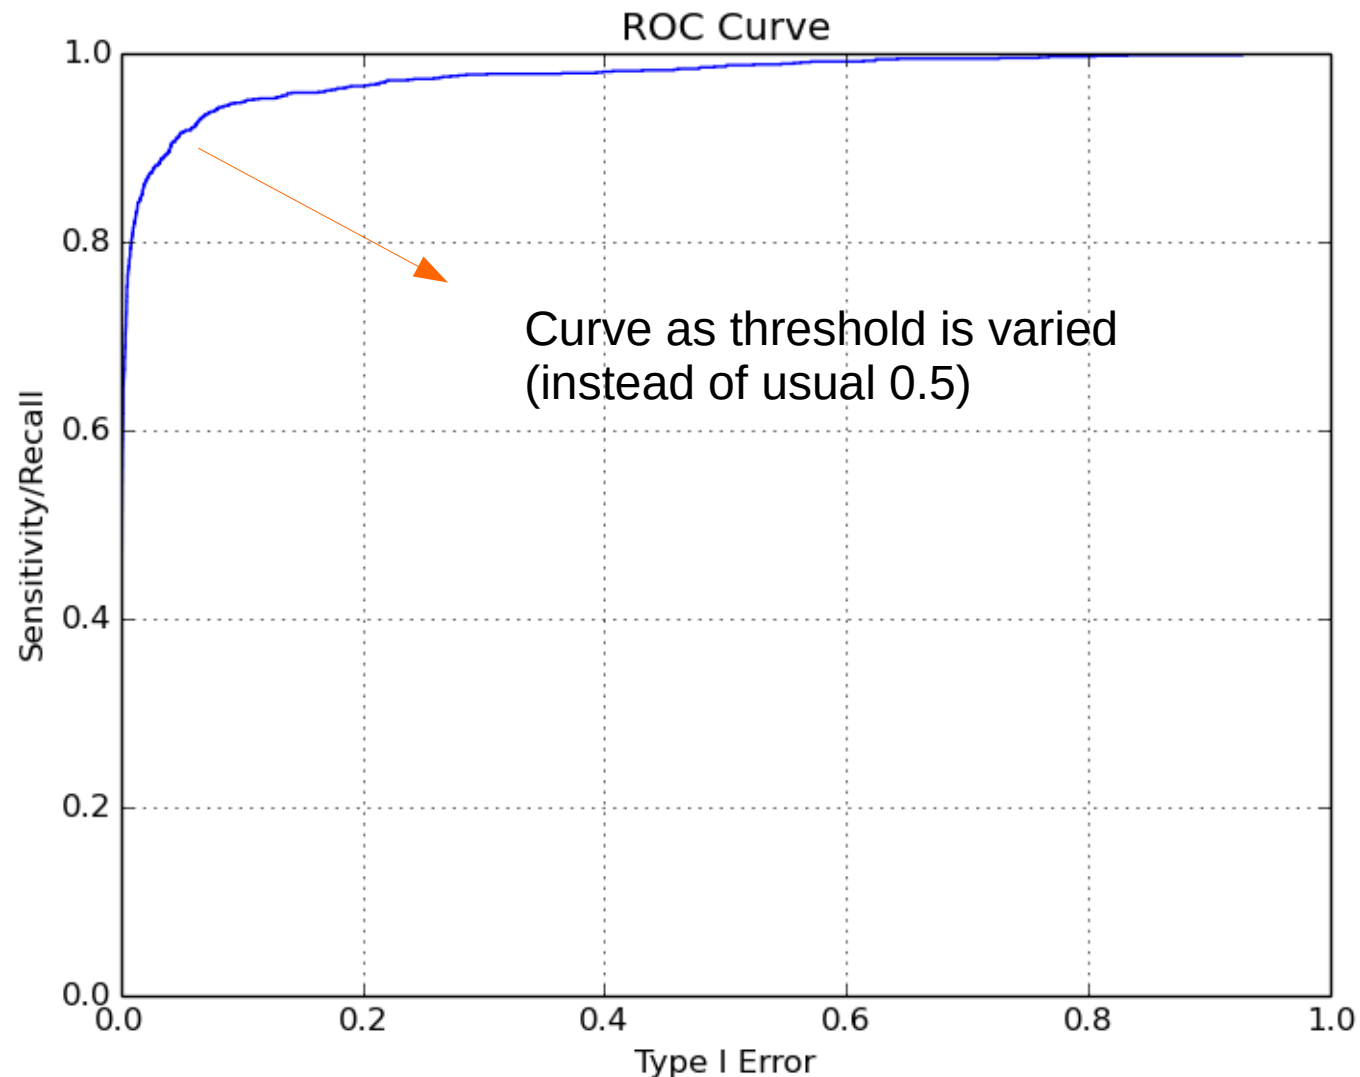
Assign every test example to label = 0:

$$\text{type II error} = \frac{1,301}{1,301 + 0} = 100\%$$

Logistic Regression:

$$\text{type II error} = \frac{224}{224 + 1,077} = 17.2\%$$

# Logistic Regression: Example 3



Area under curve,  $AUC = 0.975$

Questions?

Questions?